

Unit 6

Linear Classification & Logistic Regression

EL-GY 6143/CS-GY 6923: INTRODUCTION TO MACHINE LEARNING
PROF. PEI LIU

Learning Objectives

- ❑ Formulate a machine learning problem as a classification problem
 - Identify features, class variable, training data
- ❑ Visualize classification data using a scatter plot.
- ❑ Describe a linear classifier as an equation and on a plot.
 - Determine visually if data is perfect linearly separable.
- ❑ Formulate a classification problem using logistic regression
 - Binary and multi-class
 - Describe the logistic and soft-max function
 - Logistic function to approximate the probability
- ❑ Derive the loss function for ML estimation of the weights in logistic regression
- ❑ Use sklearn packages to fit logistic regression models
- ❑ Measure the accuracy of classification
- ❑ Adjust threshold of classifiers for trading off types of classification errors. Draw a ROC curve.
- ❑ Perform LASSO regularization for feature selection



Outline

 Motivating Example: Classifying a breast cancer test

- ☐ Linear classifiers
- ☐ Logistic regression
- ☐ Fitting logistic regression models
- ☐ Measuring accuracy in classification



NYU



TANDON SCHOOL
OF ENGINEERING

3



Diagnosing Breast Cancer

- ❑ Fine needle aspiration of suspicious lumps
- ❑ Cytopathologist visually inspects cells
 - Sample is stained and viewed under microscope
- ❑ Determines if cells are benign or malignant
 - Also provides grading if malignant
- ❑ Uses many features:
 - Size and shape of cells, degree of mitosis, differentiation, ...
- ❑ Diagnosis is not exact
- ❑ If uncertain, use a more comprehensive biopsy
 - Additional cost and time
 - Stress to patient
- ❑ Can machine learning provide better rules?



Grades of carcinoma cells
<http://breast-cancer.ca/5a-types/>



Data

- ❑ Univ. Wisconsin study, 1994
- ❑ 569 samples
- ❑ 10 visual features for each sample
- ❑ Ground truth determined by biopsy

❑ First publication: O.L. Mangasarian, W.N. Street and W.H. Wolberg. Breast cancer diagnosis and prognosis via linear programming. Operations Research, 43(4), pages 570-577, July-August 1995.

Breast Cancer Wisconsin (Diagnostic) Data Set

Download: [Data Folder](#), [Data Set Description](#)

Abstract: Diagnostic Wisconsin Breast Cancer Database



Data Set Characteristics:	Multivariate	Number of Instances:	569	Area:	Life
Attribute Characteristics:	Real	Number of Attributes:	32	Date Donated	1995-11-01
Associated Tasks:	Classification	Missing Values?	No	Number of Web Hits:	442524

Attribute Information:

- 1) ID number
- 2) Diagnosis (M = malignant, B = benign) 3-32)

Ten real-valued features are computed for each cell nucleus:

- a) radius (mean of distances from center to points on the perimeter)
- b) texture (standard deviation of gray-scale values)
- c) perimeter
- d) area
- e) smoothness (local variation in radius lengths)
- f) compactness ($\text{perimeter}^2 / \text{area} - 1.0$)
- g) concavity (severity of concave portions of the contour)
- h) concave points (number of concave portions of the contour)
- i) symmetry
- j) fractal dimension ("coastline approximation" - 1)



NYU

TANDON SCHOOL
OF ENGINEERING

Demo on Github

□ Github:

https://github.com/plugithub/MachineLearning/tree/master/unit06_logistic/breast_cancer.ipynb

Breast Cancer Diagnosis via Logistic Regression

In this demo, we will see how to visualize training data for classification, plot the logistic function and perform logistic regression. As an example, we will use the widely-used breast cancer data set. This data set is described here:

<https://archive.ics.uci.edu/ml/machine-learning-databases/breast-cancer-wisconsin>

Each sample is a collection of features that were manually recorded by a physician upon inspecting a sample of cells from fine needle aspiration. The goal is to detect if the cells are benign or malignant.

Loading and Visualizing the Data

We first load the packages as usual.

```
: import numpy as np
import matplotlib
import matplotlib.pyplot as plt
import pandas as pd
from sklearn import datasets, linear_model, preprocessing
%matplotlib inline
```

Next, we load the data. It is important to remove the missing values.

```
: names = ['id', 'thick', 'size_unif', 'shape_unif', 'marg', 'cell_size', 'bare',
          'chrom', 'normal', 'mit', 'class']
df = pd.read_csv('https://archive.ics.uci.edu/ml/machine-learning-databases/' +
                'breast-cancer-wisconsin/breast-cancer-wisconsin.data',
                names=names, na_values='?', header=None)
df = df.dropna()
df.head(6)
```

	id	thick	size_unif	shape_unif	marg	cell_size	bare	chrom	normal	mit	class
0	84	4	1	1	1	1	1	1	1	1	1
1	85	4	1	1	1	1	1	1	1	1	1
2	86	4	1	1	1	1	1	1	1	1	1
3	87	4	1	1	1	1	1	1	1	1	1
4	88	4	1	1	1	1	1	1	1	1	1



NYU

TANDON SCHOOL
OF ENGINEERING

Loading The Data

```
names = ['id','thick','size_unif','shape_unif','marg','cell_size','bare',  
         'chrom','normal','mit','class']  
df = pd.read_csv('https://archive.ics.uci.edu/ml/machine-learning-databases/' +  
                'breast-cancer-wisconsin/breast-cancer-wisconsin.data',  
                names=names, na_values='?', header=None)  
df = df.dropna()  
df.head(6)
```

	id	thick	size_unif	shape_unif	marg	cell_size	bare	chrom	normal	mit	class
0	1000025	5	1	1	1	2	1.0	3	1	1	2
1	1002945	5	4	4	5	7	10.0	3	2	1	2
2	1015425	3	1	1	1	2	2.0	3	1	1	2
3	1016277	6	8	8	1	3	4.0	3	7	1	2
4	1017023	4	1	1	3	2	1.0	3	1	1	2
5	1017122	8	10	10	8	7	10.0	9	7	1	4

- ❑ Follow standard pandas routine
- ❑ All code in Lect06_Demo.ipynb

Drops missing samples


Class = 2 => benign
Class = 4 => malignant

See following for explanation of attributes

<https://archive.ics.uci.edu/ml/machine-learning-databases/breast-cancer-wisconsin/breast-cancer-wisconsin.names>



Visualizing the Data



❑ Scatter plot of points from each class

❑ Plot not informative

- Many points overlap
- Relative frequency at each point not visible


```
y = np.array(df['class'])
xnames = ['size_unif', 'marg']
X = np.array(df[xnames])

Iben = np.where(y==2)[0]
Imal = np.where(y==4)[0]

plt.plot(X[Imal,0],X[Imal,1], 'r.')
plt.plot(X[Iben,0],X[Iben,1], 'g.')
plt.xlabel(xnames[0], fontsize=16)
plt.ylabel(xnames[1], fontsize=16)
plt.ylim(0,14)
plt.legend(['malign', 'benign'], loc='upper right')
```



Improving the Plot



❑ Make circle size proportional to count

❑ Many gymnastics to make this plot in python

```
# Compute the bin edges for the 2d histogram
x0val = np.array(list(set(X[:,0]))).astype(float)
x1val = np.array(list(set(X[:,1]))).astype(float)
x0, x1 = np.meshgrid(x0val,x1val)
x0e= np.hstack((x0val,np.max(x0val)+1))
x1e= np.hstack((x1val,np.max(x1val)+1))


# Make a plot for each class
yval = [2,4]
color = ['g','r']
for i in range(len(yval)):
    I = np.where(y==yval[i])[0]
    cnt, x0e, x1e = np.histogram2d(X[I,0],X[I,1],[x0e,x1e])
    x0, x1 = np.meshgrid(x0val,x1val)
    plt.scatter(x0.ravel(), x1.ravel(), s=2*cnt.ravel(),alpha=0.5,
               c=color[i],edgecolors='none')

plt.ylim([0,14])
plt.legend(['benign','malign'], loc='upper right')
plt.xlabel(xnames[0], fontsize=16)
plt.ylabel(xnames[1], fontsize=16)
```



In-Class Exercise

- ❑ Get into groups
 - At least one must have a laptop with jupyter notebook
- ❑ Determine a classification rule
 - Predict class label from the two features
- ❑ Test in python
 - Make the predictions
 - Measure the accuracy



In-Class Exercise

Based on the above plot, what would be a good "classifier" using the two features. That is, write a function that makes a prediction \hat{y} of the class label y . Code up your classifier function. Measure the accuracy of the classifier on the data. What percentage error does your classifier get?

```
# TODO
```




NYU

TANDON SCHOOL
OF ENGINEERING

10



A Possible Classification Rule



□ From inspection, benign if:

$$\text{marg} + \frac{2}{3}(\text{size_unif}) < 4$$

□ Classification rule from **linear constraint**

□ What are other possible classification rules?

□ Every rule misclassifies some points

□ What is optimal?



Mangasarian's Original Paper

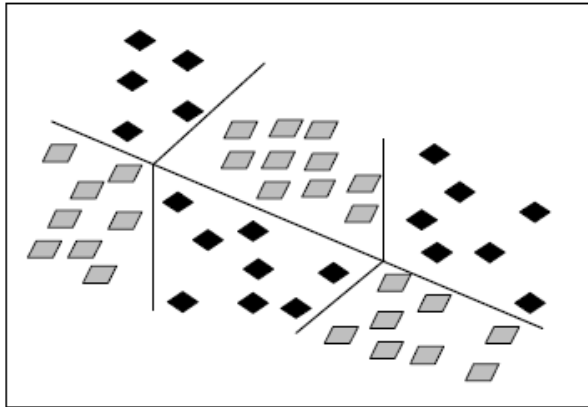


Figure 2.2 - Decision boundaries generated by MSM-T. Dark objects represent benign tumors while light object represent malignant ones.

- ❑ Proposes Multisurface method – Tree (MSM-T)
 - Decision tree based on linear rules in each step
- ❑ Fig to left from
 - Pantel, “Breast Cancer Diagnosis and Prognosis,” 1995
- ❑ Best methods today use neural networks
- ❑ This lecture will look at **linear classifiers**
 - These are much simpler
 - Do not provide same level of accuracy
- ❑ But, building block to more complex classifiers




NYU

TANDON SCHOOL
OF ENGINEERING

Outline

- ❑ Motivating Example: Classifying a breast cancer test



- ❑ Linear classifiers

- ❑ Logistic regression

- ❑ Fitting logistic regression models

- ❑ Measuring accuracy in classification



NYU

TANDON SCHOOL
OF ENGINEERING

Classification

□ Given features \mathbf{x} , determine its class label, $y = 1, \dots, K$

□ Binary classification: $y = 0$ or 1



□ Many applications:

- Face detection: Is a face present or not?
- Reading a digit: Is the digit 0,1,...,9?
- Are the cells cancerous or not?
- Is the email spam?

□ Equivalently, determine classification function:

$$\hat{y} = f(\mathbf{x}) \in \{1, \dots, K\}$$

- Like regression, but with a discrete response
- May index $\{1, \dots, K\}$ or $\{0, \dots, K - 1\}$



NYU

TANDON SCHOOL
OF ENGINEERING

Linear Classifier

□ General binary classification rule:

$$\hat{y} = f(x) = 0 \text{ or } 1$$


□ Linear classification rule:

- Take linear combination $z = w_0 + \sum_{j=1}^d w_d x_d$
- Predict class from z

$$\hat{y} = \begin{cases} 1 & z \geq 0 \\ 0 & z < 0 \end{cases}$$

□ Decision regions described by a **half-space**.


□ $\mathbf{w} = (w_0, \dots, w_d)$ is called the weight vector



NYU

TANDON SCHOOL
OF ENGINEERING

Breast Cancer Example



```
1 # A simple function with a linear decision rule
2 def predict(X):
3     marg = X[:,1]
4     size_unif = X[:,0]
5     z = marg + 2/3*size_unif - 4
6     yhat = (z > 0).astype(int)
7     return yhat
8
9 # Test on the data
10 yhat = predict(X)
11 acc = np.mean(y == yhat)
12 print('Accuracy = %7.4f' % acc)
```

Accuracy = 0.9268

□ From inspection, benign if:

$$\text{marg} + \frac{2}{3}(\text{size_unif}) < 4$$

□ Mathematically:

- $z = w_0 + w_1(\text{marg}) + w_2(\text{size_unif})$

- $\mathbf{w} = [-4, 1, \frac{2}{3}]$

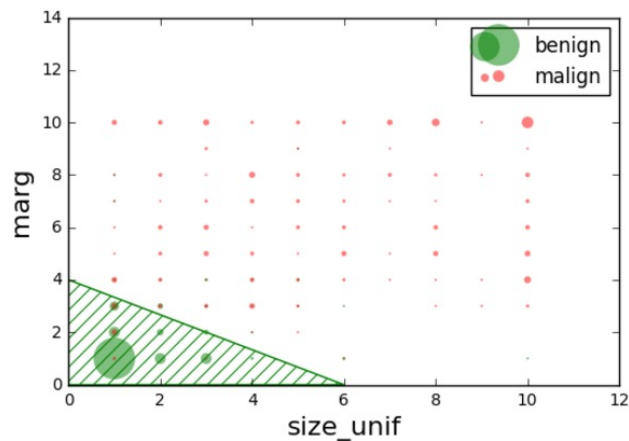
- $\hat{y} = \begin{cases} 1 & z \geq 0 \\ 0 & z < 0 \end{cases}$

□ Classification rule from **linear constraint**

- Gets 93% accuracy with just 2 features!



Breast Cancer Example



Questions for today:

- How do we use all 10 features?
- How do we fit an optimal linear classifier?
- What is optimal?



```
1 # A simple function with a linear decision rule
2 def predict(X):
3     marg = X[:,1]
4     size_unif = X[:,0]
5     z = marg + 2/3*size_unif - 4
6     yhat = (z > 0).astype(int)
7     return yhat
8
9 # Test on the data
10 yhat = predict(X)
11 acc = np.mean(y == yhat)
12 print('Accuracy = %7.4f' % acc)
```

Accuracy = 0.9268



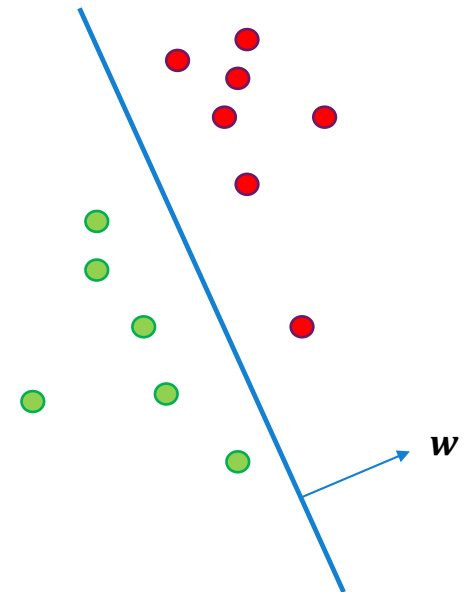
Linear vs. Non-Linear

- ❑ Linear boundaries are limited
- ❑ Can only describe very simple regions
- ❑ But, serves as building block
 - Many classifiers use linear rules as first step
 - Neural networks, decision trees, ...
- ❑ Breast cancer example:
 - Is the region linear or non-linear?



Perfect Linear Separability

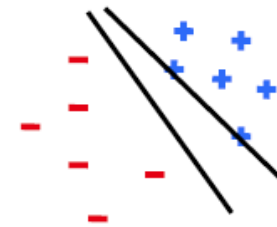
- Given training data $(\mathbf{x}_i, y_i), i = 1, \dots, N$
 - Binary class label: $y_i = \{0, 1\}$
- **Perfectly linearly separable** if:
there is a linear classifier that makes no errors on the training data
- Visually: You can draw a line between the points
- Mathematically: There exists a $\mathbf{w} = (w_0, w_1, \dots, w_d)$ such that:
 - $w_0 + w_1 x_{i1} + \dots + w_d x_{id} > 0$ when $y_i = 1$
 - $w_0 + w_1 x_{i1} + \dots + w_d x_{id} < 0$ when $y_i = 0$



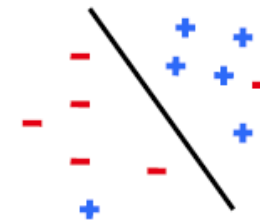
Most Data not Perfectly Separable

- ❑ Generally cannot find a separating hyperplane
- ❑ Always, some points that will be mis-classified
- ❑ Algorithms attempt to find “good” hyper-planes
 - Reduce the number of mis-classified points
 - Or, some similar metric

Separable




Non-Separable



Non-Uniqueness

- When one exists, separating hyper-plane is not unique
- Fig. on right: Many separating planes
- Example:
 - If \mathbf{w} is separating, then so is $\alpha \mathbf{w}$ for all $\alpha > 0$
- Which one is optimal?



NYU

TANDON SCHOOL
OF ENGINEERING

In-Class Exercise

Question 1. Linear Classifier

Given the data below, use the `plt.scatter()` function to plot the data points with different colors for the classes.

```
X = np.array([[1,1], [1,3], [2,2], [2,3], [3,2], [4,1]])
y = np.array([0,0,1,1,0,1], dtype=np.int)

# TODO
```

You should see that the data is not linearly separable. Find a linear classifier that makes a minimal number of errors on the training data.

Write a function `predict()` function for the classifier and get the predicted labels with the command:

```
yhat = predict(X)
```

Print `yhat` and `y`. How many errors does your classifier make?

- Complete in `logistic_inclass.ipynb`
 - Can use Google colab or your local machine



Outline

- ❑ Motivating Example: Classifying a breast cancer test

- ❑ Linear classifiers

- ❑ Logistic regression

- ❑ Fitting logistic regression models

- ❑ Measuring accuracy in classification



NYU

TANDON SCHOOL
OF ENGINEERING

23



Hard vs. Soft Decision Classifiers

❑ Binary classification problem:

- Given features x , estimate class label 0 or 1
- Ex: cat vs. dog

❑ Hard decision classifier:

- Output a class label: $\hat{y} = 0$ or 1
- Ex: $\hat{y} = 1 \Rightarrow \text{Image is a dog!}$

❑ Soft decision classifier:

- Output a **conditional probability** $P(y = 1|x)$
- $P(y = 1|x)$ is between 0 and 1
- Ex: $P(y = 1|x) = 0.9 \Rightarrow \text{Given this image, there is a 90\% chance it is a dog}$

Cat
 $y = 0$

Dog
 $y = 1$



NYU

TANDON SCHOOL
OF ENGINEERING

24



Why Use Soft Decision Classifiers?

- ❑ In most problems, classifiers make errors
- ❑ Example: Is the digit a 5 or 6?
 - Hard decision makes decision with certainty
 - But the decision can be wrong
- ❑ Soft decision classifiers recognize this uncertainty
- ❑ Easier to train soft decision classifier
 - Allows for error in training data
 - See next section
- ❑ Provides a confidence measure

Example from MNIST dataset
True digit = 5



Hard decision
Digit = 6!



Soft decision
Digit = 5 with probability 30%
Digit = 6 with probability 70%




Logistic Model for Binary Classification

❑ Binary classification problem: $y = 0, 1$


❑ Hard decision linear classifier

- Predict a class label $\hat{y} = 0$ or 1
- $z = w_0 + \sum_j w_j x_j$
- $\hat{y} = \begin{cases} 1 & z > 0 \\ 0 & z \leq 0 \end{cases}$




❑ Logistic soft decision classifier

- Predict a probability $P(y = 1|x)$
- $z = w_0 + \sum_j w_j x_j$
- $P(y = 1|x) = \frac{1}{1+e^{-z}}$



Logistic Model as a “Soft” Classifier



□ Plot of

$$P(y = 1|x) = \frac{1}{1 + e^{-z}}, \quad z = w_1 x$$

◦ Markers are random samples

□ Higher w_1 : prob transition becomes sharper

◦ Fewer samples occur across boundary

□ As $w_1 \rightarrow \infty$ logistic becomes “hard” rule

$$P(y = 1|x) \approx \begin{cases} 1 & x > 0 \\ 0 & x < 0 \end{cases}$$



NYU

TANDON SCHOOL
OF ENGINEERING


Hard vs. Soft Classification in 2D

□ Hard decision:

- Divides space into two half-space
- One for each class label

□ Soft decision

- Gradual transition for probability 0 to 1



Multi-Class Logistic Regression

- ❑ Logistic regression easily extended to multiple classes
- ❑ Suppose $y \in 1, \dots, K$
 - K possible classes (e.g. digits, letters, spoken words, ...)
- ❑ Two parameters: $\mathbf{W} \in R^{K \times d}$, $\mathbf{w}_0 \in R^K$ Slope matrix and bias
- ❑ Step 1: Create K linear functions.

$$\mathbf{z} = \mathbf{W}\mathbf{x} + \mathbf{w}_0$$

- Called **scores** or **logits**

- ❑ Step 2: Predict probabilities via **softmax function**

$$P(y = k|\mathbf{x}) = \frac{e^{z_k}}{\sum_{\ell=1}^K e^{z_\ell}}$$



Softmax Example

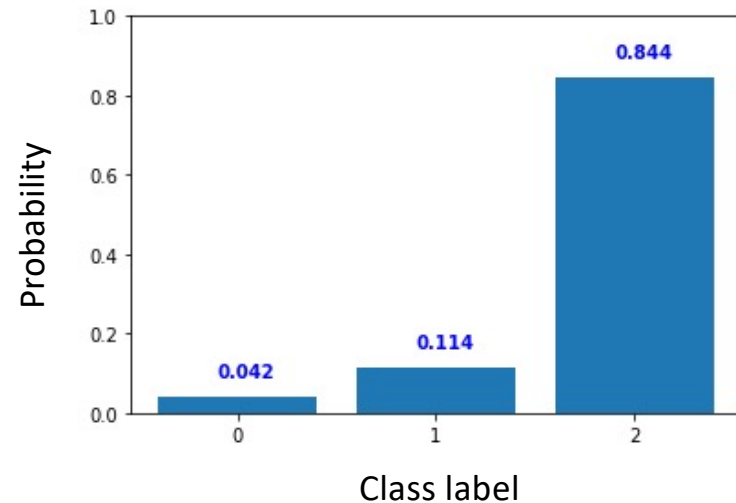
□ Suppose:

- Three class $y \in \{0,1,2\}$
- For some x , the logits are $z = [-1, 0, 2]$

□ Then $\exp(z) = [0.36, 1, 7.36]$

□ Softmax probabilities:

- $P(y = 0|x) = \frac{0.36}{0.36+1+7.36} \approx 0.042$
- $P(y = 1|x) = \frac{1}{0.36+1+7.36} \approx 0.114$
- $P(y = 2|x) = \frac{7.36}{0.36+1+7.36} \approx 0.844$



Softmax Properties

□ Consider **soft-max** function:

$$g_k(\mathbf{z}) = \frac{e^{z_k}}{\sum_{\ell=1}^K e^{z_\ell}}$$

- K inputs $\mathbf{z} = (z_1, \dots, z_K)$
- K outputs $g(\mathbf{z}) = (g(\mathbf{z})_1, \dots, g(\mathbf{z})_K)$


□ $g(\mathbf{z})$ is like a PMF on the labels $[0, 1, \dots, K - 1]$

- $g_k(\mathbf{z}) \in [0, 1]$ for each component k
- $\sum_{k=1}^K g_k(\mathbf{z}) = 1$


□ Softmax property: When $z_k \gg z_\ell$ for all $\ell \neq k$:

- $g_k(\mathbf{z}) \approx 1$
- $g_\ell(\mathbf{z}) \approx 0$ for all $\ell \neq k$

□ Assigns highest probability to class k when z_k is largest



Multi-Class Logistic Regression Decision Regions



- ❑ Each decision region defined by set of hyperplanes
- ❑ Intersection of linear constraints
- ❑ Sometimes called a **polytope**



Transform Linear Models

□ As in regression, logistic models can be applied to transform features

□ **Step 1:** Map \mathbf{x} to some transform features, $\phi(\mathbf{x}) = [\phi_1(\mathbf{x}), \dots, \phi_p(\mathbf{x})]^T$ ← Additional transform step

□ **Step 2:** Linear weights: $z_k = \sum_{j=1}^p W_{kj} \phi_j(\mathbf{x})$

□ **Step 3:** Soft-max $P(y = k | \mathbf{z}) = g_k(\mathbf{z}), \quad g_k(\mathbf{z}) = \frac{e^{z_k}}{\sum_{\ell} e^{z_{\ell}}}$


□ **Example transforms:**

- Standard regression $\phi(\mathbf{x}) = [1, x_1, \dots, x_j]^T$ (j original features, $j+1$ transformed features)
- Polynomial regression: $\phi(\mathbf{x}) = [1, x, \dots, x^d]^T$ (1 original feature, $d + 1$ transformed features)



Using Transformed Features

- ❑ Enables richer class boundaries
- ❑ Example: Fig B is not linearly separable
- ❑ But, consider nonlinear features
 - $\phi(x) = [1, x_1, x_2, x_1^2, x_2^2]^T$
- ❑ Then can discriminate classes with linear function
 - $w = [-r^2, 0, 0, 1, 1]$
 - $z = w^T \phi(x) = x_1^2 + x_2^2 - r^2$
 - Blue when $z \leq 0$ and Green when $z > 0$



In-Class Exercise

Question 2: Logistic Model

Consider the model for the passing a test:

$$P(\text{pass test}) = 1/(1+\exp(-z)), \quad z = w_0 + w_1 \cdot \text{hrs_alone} + w_2 \cdot \text{hrs_tutor}$$

where `hrs_alone` is the number of hours studied alone and `hrs_tutor` is the number of hours with a tutor. Given the values below find w_0 for the probability = 0.6.

```
: hrs_alone = 4
  hrs_tutor = 1
  w1 = 0.2
  w2 = 0.5
  prob = 0.6

# TODO
# w0 = ...
```

Given the values above, plot the probability of passing as a function of `hrs_tutor` in the range of 0 to 10 hours.

- ❑ Complete in `logistic_inclass.ipynb`
 - Can use Google colab or your local machine



Outline

- ❑ Motivating Example: Classifying a breast cancer test

- ❑ Linear classifiers

- ❑ Logistic regression

- ❑ Fitting logistic regression models

- ❑ Measuring accuracy in classification



NYU

TANDON SCHOOL
OF ENGINEERING

36



How To Fit Logistic Models?


- Given training data, $(\mathbf{x}_i, y_i), i = 1, \dots, N$
 - Binary labels $y_i \in \{0,1\}$

- **Binary logistic model:** Given *new* \mathbf{x} predict class probability via:

- Linear weights: $\mathbf{z} = \mathbf{w}_{1:p}^T \mathbf{x} + w_0$
- Sigmoid: $P(y = 1|\mathbf{x}) = \frac{1}{1+e^{-\mathbf{z}}}$

- Weight vector \mathbf{w} represents unknown **model parameters**

- **Learning problem:** Learn weight vector \mathbf{w} from the data



Maximum Likelihood Principle

□ **Likelihood function:** From the logistic model, we can derive:

$P(\mathbf{y}|\mathbf{X}, \mathbf{w})$ = Probability of class labels given inputs \mathbf{X} and weights \mathbf{w}

- (\mathbf{X}, \mathbf{y}) are the data matrices for all n training samples
- \mathbf{w} is the vector of parameters

□ **Key idea:** $P(\mathbf{y}|\mathbf{X}, \mathbf{w})$ is higher \Rightarrow data is a better match with the parameters

□ **Maximum Likelihood Principle:** Given data (\mathbf{X}, \mathbf{y}) :

Find parameters \mathbf{W} to maximize $P(\mathbf{y}|\mathbf{X}, \mathbf{W})$



Binary Cross Entropy

□ Given data $(x_i, y_i), i = 1, \dots, N$ with binary labels $y_i \in \{0,1\}$

□ **Theorem:** MLE for logistic model is equivalent to minimizing the **binary cross entropy**:

$$J(\mathbf{w}) = \sum_{i=1}^n \ln[1 + e^{z_i}] - y_i z_i, \quad z_i = w_0 + \sum_{j=1}^d w_j x_{ij}$$

- Find the weight vector \mathbf{w} to minimize $J(\mathbf{w})$
- Will prove below this is equivalent to maximizing $P(\mathbf{y}|\mathbf{X}, \mathbf{w})$

□ Provides a simple cost function to minimize for fitting

□ Note that z_i are implicitly function of weights \mathbf{w}



Visualizing BCE


Binary cross entropy

$$J(\mathbf{w}) = \sum_{i=1}^n \ln[1 + e^{z_i}] - y_i z_i, \quad z_i = w_0 + \sum_{j=1}^d w_j x_{ij}$$

Each term has cost $\ln[1 + e^{z_i}] - y_i z_i$

$y_i = 0 \Rightarrow$ Tries to make z_i negative

$y_i = 1 \Rightarrow$ Tries to make z_i positive



Higher cost for $y = 1$

Higher cost for $y = 0$



Min and Argmin

□ Given a function $f(x)$

□ $\min_x f(x)$


- Minimum value of the $f(x)$
- Point on the y -axis

□ $\arg \min_x f(x)$

- Value of x where $f(x)$ is a minimum
- Point on the x -axis

□ Similarly, define $\max_x f(x)$ and $\arg \max_x f(x)$

$$f(x) = (x - 1)^2 + 2$$



$$\min_x f(x) = 2$$

$$\arg \min_x f(x) = 1$$



MLE Using Argmax

□ We can write the MLE using argmax

□ Suppose we have

- Data (\mathbf{X}, \mathbf{y})
- Likelihood function $P(\mathbf{y}|\mathbf{X}, \mathbf{w})$

□ Then, MLE is equivalent to:

$$\hat{\mathbf{w}} = \arg \max_{\mathbf{w}} P(\mathbf{y}|\mathbf{X}, \mathbf{w})$$

- Equivalent to saying find \mathbf{w} to maximize $P(\mathbf{y}|\mathbf{X}, \mathbf{w})$



Log Likelihood

□ Assume outputs y_i are independent, depending only on x_i

□ Then, likelihood factors:

$$P(\mathbf{y}|\mathbf{X}, \mathbf{w}) = \prod_{i=1}^N P(y_i|x_i, \mathbf{w})$$


□ Define **negative log likelihood**:

$$L(\mathbf{w}) = -\ln P(\mathbf{y}|\mathbf{X}, \mathbf{w}) = -\sum_{i=1}^N \ln P(y_i|x_i, \mathbf{w})$$

□ Maximum likelihood estimator can be re-written as:

$$\hat{\mathbf{w}} = \arg \max_{\mathbf{w}} P(\mathbf{y}|\mathbf{X}, \mathbf{w}) = \arg \min_{\mathbf{w}} L(\mathbf{w})$$

Negative Log Likelihood



NYU

TANDON SCHOOL
OF ENGINEERING

Cross Entropy

- The cross-entropy of the distribution q relative to a distribution p is defined as follows:

$$H(p, q) = -\sum_x p(x) \log(q(x))$$

- Similar to entropy of a distribution:

$$H(p) = -\sum_x p(x) \log(p(x))$$

- It measures the difference in distribution p and q
 - Cross-entropy is minimum when two distributions are equal
 - Otherwise, it is $H(p)$ + Divergence between p and q (Kullback–Leibler)



Logistic Loss = Binary Cross Entropy

□ Negative log likelihood function: $J(\mathbf{w}) = -\sum_{i=1}^n \ln P(y_i|\mathbf{x}_i, \mathbf{w})$

$$P(y_i = 1|\mathbf{x}_i, \mathbf{w}) = \frac{1}{1 + e^{-z_i}}, \quad z_i = \mathbf{w}_{1:p}^T \mathbf{x}_i + w_0$$

□ Therefore,

$$P(y_i = 1|\mathbf{x}_i, \mathbf{w}) = \frac{e^{z_i}}{1 + e^{z_i}}, \quad P(y_i = 0|\mathbf{x}_i, \mathbf{w}) = \frac{1}{1 + e^{z_i}}$$

□ Hence,

$$\ln P(y_i|\mathbf{x}_i, \mathbf{w}) = y_i \ln P(y_i = 1|\mathbf{x}_i, \mathbf{w}) + (1 - y_i) \ln P(y_i = 0|\mathbf{x}_i, \mathbf{w}) = y_i z_i - \ln[1 + e^{z_i}]$$

□ Loss function = binary cross entropy:

$$J(\mathbf{w}) = \sum_{i=1}^n \ln[1 + e^{z_i}] - y_i z_i$$



Multi-Class Classification

□ For multi-class classification, define the “one-hot” vector:

$$r_{ik} = \begin{cases} 1 & y_i = k \\ 0 & y_i \neq k \end{cases}, \quad i = 1, \dots, N, \quad k = 1, \dots, K$$

□ Then, $\ln P(y_i | \mathbf{x}_i, \mathbf{W}) = \sum_{k=1}^K r_{ik} \ln P(y_i = k | \mathbf{x}_i, \mathbf{W})$

□ Hence, negative log likelihood is:

$$J(\mathbf{W}) = \sum_{i=1}^N \left[\ln \left[\sum_k e^{z_{ik}} \right] - \sum_k z_{ik} r_{ik} \right]$$

- Sometimes called the **categorical cross-entropy**
- Can prove this with some algebra



Gradient Calculations

- ❑ To minimize take partial derivatives: $\frac{\partial J(W)}{\partial W_{kj}} = 0$ for all W_{kj}
- ❑ Define transform matrix: $A_{ij} = \phi_j(\mathbf{x}_i)$
- ❑ Hence, $z_{ik} = \sum_{j=1}^p A_{ij} W_{kj}$
- ❑ Estimated class probabilities: $p_{ik} = \frac{e^{z_{ik}}}{\sum_{\ell} e^{z_{i\ell}}}$
- ❑ Gradient components are: $\frac{\partial L(W)}{\partial W_{kj}} = \sum_{i=1}^N (p_{ik} - r_{ik}) A_{ij} = 0$
 - $K \times p$ equations and $K \times p$ unknowns
- ❑ Unfortunately, no closed-form solution to these equations
 - Nonlinear dependence of p_{ik} on terms in W



Numerical Optimization

□ We saw that we can find minima by setting $\nabla f(x) = 0$

- M equations and M unknowns.
- May not have closed-form solution

□ **Numerical methods:** Finds a sequence of estimates x^t
 $x^t \rightarrow x^*$

- Under some conditions, it converges to some other “good” minima
- Run on a computer program, like python

□ Next unit: Will discuss numerical methods to perform optimization

□ This lecture: Use built-in python routine



NYU

TANDON SCHOOL
OF ENGINEERING

48



In-Class Exercise

Question 3. Calculating and Plotting the Binary Cross Entropy Loss

You are given the scalar data x and y with binary class labels below.

```
x = np.array([-1,1,3,4,5])
y = np.array([0,0,1,0,1])
```

```
def bce_loss(x,y,w):
    # TODO
    # J = BCE Loss
    return J

# Print the loss for `w_manual`
```

Now consider a set of $w = [w_0, 0.5]$.

- Plot the BCE loss over 100 values w_0 from -2.5 to 0
- What value of w_0 gives the minimum BCE loss? Call this w_{0_opt} .
- What is the minimum BCE loss?

- Complete in `logistic_inclass.ipynb`
 - Can use Google colab or your local machine




NYU

TANDON SCHOOL
OF ENGINEERING

49



Outline

- ❑ Motivating Example: Classifying a breast cancer test
- ❑ Linear classifiers
- ❑ Logistic regression
- ❑ Fitting logistic regression models
-  ❑ Returning to the breast cancer dataset
- ❑ Measuring accuracy in classification



NYU

TANDON SCHOOL
OF ENGINEERING

50



Fitting Two Variables

```
1 xnames = ['size_unif', 'marg']
2 X = np.array(df[xnames])
3 print(X.shape)
```

(683, 2)

```
scal = StandardScaler()
Xtr1 = scal.fit_transform(Xtr)
Xts1 = scal.transform(Xts)
```

```
reg = linear_model.LogisticRegression(C=1e5)
reg.fit(Xtr1, ytr)
```

```
1 yhat = reg.predict(Xts1)
2 acc = np.mean(yhat == yts)
3 print("Accuracy on test data = %f" % acc)
```

Accuracy on test data = 0.917073

❑ Get the two variables

❑ Scale, fit and test

❑ Regularization note for sklearn


- Always performs regularized regression:
- Minimizes

$$J(w) + \frac{1}{2C} \|w\|^2$$

- $J(w)$ = Negative log likelihood
- $\frac{1}{2C} \|w\|^2$ = regularization term (like Ridge)
- Used for numerical stability



Visualizing the Decision Regions



Probability output:

- Smooth value from 0 to 1

Class decision:

- $\hat{y} = 1$ when $P(y = 1|x) > 0.5$
- $\hat{y} = 0$ when $P(y = 1|x) \leq 0.5$
- Decision boundary is a line
- Similar to what we manually selected



Fitting All The Variables

```
1 # Get array of all the features. These are the columns in the dataframe
2 # except the last
3 xnames = names[:-1]
4 X = np.array(df[xnames])
5 print(X.shape)
6
7 # Split into training and test
8 Xtr, Xts, ytr, yts = train_test_split(X, y, test_size=0.30)
```

```
1 # Scale the data
2 scal = StandardScaler()
3 Xtr1 = scal.fit_transform(Xtr)
4 Xts1 = scal.transform(Xts)
5
6 # Fit on the scaled trained data
7 reg = linear_model.LogisticRegression(C=1e5)
8 reg.fit(Xtr1, ytr)
9
10 # Measure accuracy
11 yhat = reg.predict(Xts1)
12 acc = np.mean(yhat == yts)
13 print("Accuracy on training data = %f" % acc)
```

Accuracy on training data = 0.970732

❑ With all 10 variables:

- Get 97% test accuracy

❑ 10-fold cross validation

- Also in demo
- Accuracy = 0.967
- SE = 0.011

	feature	slope
0	id	-0.513702
1	thick	1.498450
2	size_unif	-0.293459
3	shape_unif	0.702795
4	marg	1.207218
5	cell_size	0.145706
6	bare	1.153128
7	chrom	1.459208
8	normal	0.719531
9	mit	0.785216



In-Class Exercise

Question 4. Heart Attack Fit

In this exercise, we fit heart attack data from the UCI website. We can load it as follows.

```
: ## Generate synthetic data
names = ['age', 'sex', 'cp', 'trestbps', 'chol', 'fbs',
         'restecg', 'thalach', 'exang', 'oldpeak', 'slope', 'ca', 'thal', 'num']
url = 'https://archive.ics.uci.edu/ml/machine-learning-databases/heart-disease/processed.cleveland.data'
df = pd.read_csv(url, na_values='?', header=None, names=names)
df = df.dropna()
```

Print the first few rows of the data frame. Print the number of attributes of number of samples

- ❑ Complete in logistic_inclass.ipynb
 - Can use Google colab or your local machine



NYU

TANDON SCHOOL
OF ENGINEERING

54




Outline

- ❑ Motivating Example: Classifying a breast cancer test

- ❑ Linear classifiers

- ❑ Logistic regression

- ❑ Fitting logistic regression models

- ❑ Measuring accuracy in classification



NYU

TANDON SCHOOL
OF ENGINEERING

55



Errors in Binary Classification

Two types of errors:

- Type I error (False positive / false alarm): Decide $\hat{y} = 1$ when $y = 0$
- Type II error (False negative / missed detection): Decide $\hat{y} = 0$ when $y = 1$

Implication of these errors may be different

- Think of breast cancer diagnosis

Accuracy of classifier can be measured by:

- $TPR = P(\hat{y} = 1|y = 1)$
- $FPR = P(\hat{y} = 1|y = 0)$
- $Accuracy = P(\hat{y} = 1|y = 1) + P(\hat{y} = 0|y = 0)$
 - (percentage of correct classification)

predicted→ real↓	<i>Class_pos</i>	<i>Class_neg</i>
<i>Class_pos</i>	TP	FN
<i>Class_neg</i>	FP	TN

$$TPR \text{ (sensitivity)} = \frac{TP}{TP + FN}$$

$$FPR \text{ (1-specificity)} = \frac{FP}{TN + FP}$$



Many Other Metrics

From previous slide

- $TPR = P(\hat{y} = 1|y = 1)$ =sensitivity
- $FPR = P(\hat{y} = 1|y = 0)$ =1-specificity

predicted→ real↓	Class_pos	Class_neg
Class_pos	TP	FN
Class_neg	FP	TN

$$TPR \text{ (sensitivity)} = \frac{TP}{TP + FN}$$

$$FPR \text{ (1-specificity)} = \frac{FP}{TN + FP}$$

Machine learning often uses (positive=items of interests in retrieval applications)

- Recall = Sensitivity = $TP/(TP+FN)$ (How many positives are detected among all positive?)
- Precision = $TP/(TP+FP)$ (How many detected positive is actually positive?)
- F1-score = $\frac{Precision * Recall}{(Precision+Recall)/2} = \frac{2TP}{2TP+FN+FP} = \frac{TP}{TP+\frac{FN+FP}{2}}$
- Accuracy = $(TP+TF)/(TP+FP+TN+FN)$ (percentage of correct classification)

Medical tests:

- Sensitivity = $P(\hat{y} = 1|y = 1) = TPR$
- Specificity = $P(\hat{y} = 0|y = 0) = 1 - FPR$ = True negative rate
- Need a good tradeoff between sensitivity and specificity



Breast Cancer

- ❑ Measure accuracy on test data
- ❑ Use 4-fold cross-validation
- ❑ Sklearn has built-in functions for CV

```
Precision = 0.9614
Recall =    0.9554
f1 =       0.9578
Accuracy = 0.9664
```

```
from sklearn.model_selection import KFold
from sklearn.metrics import precision_recall_fscore_support
nfold = 4
kf = KFold(n_splits=nfold)
prec = []
rec = []
f1 = []
acc = []
for train, test in kf.split(Xs):
    # Get training and test data
    Xtr = Xs[train,:]
    ytr = y[train]
    Xts = Xs[test,:]
    yts = y[test]

    # Fit a model
    logreg.fit(Xtr, ytr)
    yhat = logreg.predict(Xts)

    # Measure
    preci, reci, f1i, _ = precision_recall_fscore_support(yts, yhat, average='binary')
    prec.append(preci)
    rec.append(reci)
    f1.append(f1i)
    acci = np.mean(yhat == yts)
    acc.append(acci)

# Take average values of the metrics
precm = np.mean(prec)
recm = np.mean(rec)
f1m = np.mean(f1)
accm = np.mean(acc)

print('Precision = {0:.4f}'.format(precm))
print('Recall = {0:.4f}'.format(recm))
print('f1 = {0:.4f}'.format(f1m))
print('Accuracy = {0:.4f}'.format(accm))
```



Go through the demo

- ❑ Up to binary classification with cross validation



NYU

TANDON SCHOOL
OF ENGINEERING

59



Hard Decisions

❑ Logistic classifier outputs a **soft** label: $P(y = 1|x) \in [0,1]$

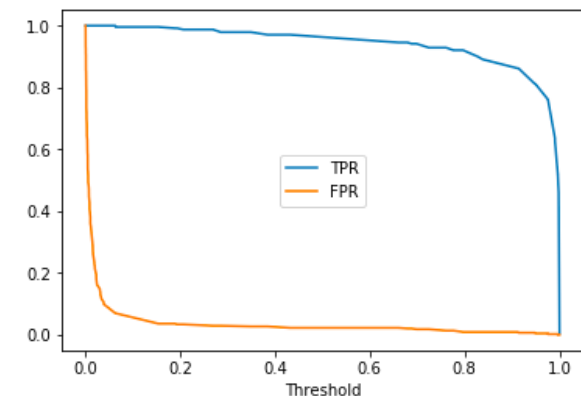
- $P(y = 1|x) \approx 1 \Rightarrow y = 1$ more likely
- $P(y = 0|x) \approx 1 \Rightarrow y = 0$ more likely

❑ Can obtain a **hard label** by **thresholding**:

- Set $\hat{y} = 1$ if $P(y = 1|x) > t$
- t = Threshold

❑ How to set threshold?

- Set $t = \frac{1}{2} \Rightarrow$ Minimizes overall error rate
- Increasing $t \Rightarrow$ Decreases false positives, but also reduces sensitivity
- Decreasing $t \Rightarrow$ Increases sensitivity, but also increases false positive





ROC Curve

- ❑ Varying threshold obtains a set of classifier
- ❑ Trades off FPR (1-specificity) and TPR (sensitivity)
- ❑ Can visualize with ROC curve
 - Receiver operating curve
 - Term from digital communications

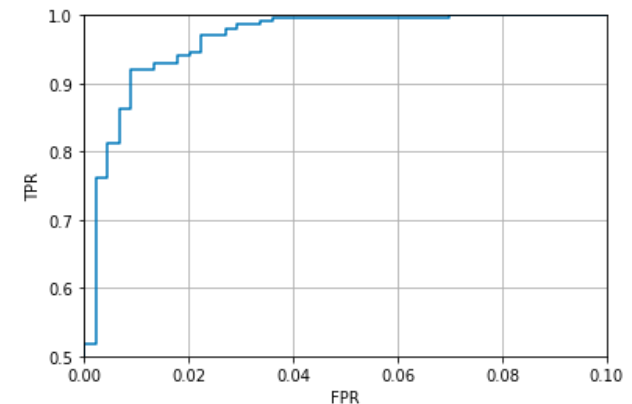
```
from sklearn import metrics
yprob = logreg.predict_proba(Xs)
fpr, tpr, thresholds = metrics.roc_curve(y, yprob[:,1])

plt.plot(fpr, tpr)
plt.grid()
plt.xlabel('FPR')
plt.ylabel('TPR')
```



Area Under the Curve (AUC)

- ❑ As one may choose a particular threshold based on the desired trade-off between the TPR and FPR, it may not be appropriate to evaluate the performance of a classifier for a fixed threshold.
- ❑ AUC is a measure of goodness for a classifier that is independent of the threshold.
- ❑ A method with a higher AUC means that under the same FPR, it has higher PPR.
- ❑ What is the highest AUC?
- ❑ Should report average AUC over cross validation folds



```
uac=metrics.roc_auc_score(y,yprob[:,1])  
print("UAC=%f" % uac)
```

UAC=0.996315



Multi-Class Classification in Python

❑ Two options

❑ One vs Rest (OVR)

- Solve a binary classification problem for each class k
- For each class k , train on modified binary labels (indicates if sample is in class or not)

$$\tilde{y}_i = \begin{cases} 1 & \text{if } y_i = k \\ 0 & \text{if } y_i \neq k \end{cases}$$

- Predict based on classifier that yields highest score

❑ Multinomial

- Directly solve weights for all classes using the multi-class cross entropy



NYU

TANDON SCHOOL
OF ENGINEERING

63



Metrics for Multiclass Classification

- ❑ Using a $K \times K$ confusion matrix
- ❑ Should normalize the matrix:
 - Sum over each row = 1
- ❑ Can compute accuracy:
 - Per class: This is the diagonal entry
 - Average: The average of the diagonal entries

Pred-->	1	2	...	K
Real↓				
1				
2				
...				
K				



NYU

TANDON SCHOOL
OF ENGINEERING

64



LASSO Regularization for Logistic Regression

- ❑ Similar to linear regression, we can use LASSO regularization with logistic regression
 - Forces the weighting coefficients to be sparse.

- ❑ Add L1 penalty $L(\mathbf{W}) = \sum_{i=1}^N [\ln[\sum_k e^{z_{ik}}] - z_{ik}r_{ik}] + \lambda \|\mathbf{W}\|_1$

- ❑ The regularization level λ should be chosen via cross validation as before

- ❑ Sklearn implementation:

```
logreg = linear_model.LogisticRegression(penalty='l1')
logreg.C = c
```

- Default use l2 penalty, to reduce the magnitude of weights
- ❑ C is the inverse of regularization strength ($C = 1/\lambda$); must be a positive float.
 - Should use a large C if you do not want to apply regularization
- ❑ Go through the LASSO part of the demo

Go through the demo

- Go through the last part with LASSO regression



NYU

TANDON SCHOOL
OF ENGINEERING

66



What You Should Know

- ❑ Formulate a machine learning problem as a classification problem
 - Identify features, class variable, training data
- ❑ Visualize classification data using a scatter plot.
- ❑ Describe a linear classifier as an equation and on a plot.
 - Determine visually if data is perfect linearly separable.
- ❑ Formulate a classification problem using logistic regression
 - Binary and multi-class
 - Describe the logistic and soft-max function
 - Understand the idea of using the logistic function to approximate the probability
- ❑ Derive the loss function for ML estimation of the weights in logistic regression
- ❑ Use sklearn packages to fit logistic regression models
- ❑ Measure the accuracy of classification: precision, recall, accuracy
- ❑ Adjust threshold of classifiers for trading off types of classification errors. Draw a ROC curve and determine AUC
- ❑ Perform LASSO regularization for feature selection

