# Unit 8 Support Vector Machines

EL-GY 6143/CS-GY 6923: INTRODUCTION TO MACHINE LEARNING PROF. PEI LIU





# **Learning Objectives**

- □ Interpret weights in linear classification of images
- ☐ Describe why linear classification for images does not work
- ☐ Define the margin in linear classification
- Describe the SVM classification problem.
- ☐ Write equations for solutions of constrained optimization using the Lagrangian.
- ☐ Describe a kernel SVM problem for non-linear classification
- ☐ Implement SVM classifiers in python
- ☐ Select SVM parameters from cross-validation



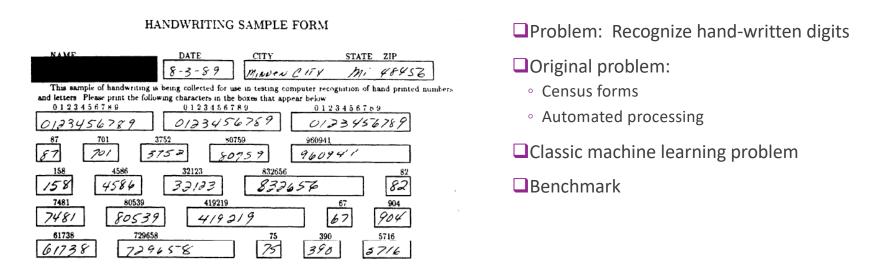


## Outline

- Motivating example: Recognizing handwritten digits
- Why logistic regression doesn't work well.
- ☐ Maximum margin classifiers
- ■Support vector machines
- ☐ Kernel trick
- lueConstrained optimization



# MNIST Digit Classification



From Patrick J. Grother, NIST Special Database, 1995





# A Widely-Used Benchmark

■We will look at SVM today

■Not the best algorithm

☐But quite good

☐...and illustrates the main points

#### Classifiers [edit]

This is a table of some of the machine learning methods used on the database and their error rates, by type of classifier:

Type \$	Classifier +	Distortion +	Preprocessing +	Error rate (%) \$	
Linear classifier	Pairwise linear classifier	None	Deskewing	7.6 <sup>[9]</sup>	
K-Nearest Neighbors	K-NN with non-linear deformation (P2DHMDM)	None	Shiftable edges	0.52 <sup>[14]</sup>	
Boosted Stumps	Product of stumps on Haar features	None	Haar features	0.87 <sup>[15]</sup>	
Non-Linear Classifier	40 PCA + quadratic classifier	None	None	3.3 <sup>[9]</sup>	1
Support vector machine	Virtual SVM, deg-9 poly, 2-pixel jittered	None	Deskewing	0.56 <sup>[16]</sup>	
Neural network	2-layer 784-800-10	None	None	1.6 <sup>[17]</sup>	1
Neural network	2-layer 784-800-10	elastic distortions	None	0.7 <sup>[17]</sup>	
Deep neural network	6-layer 784-2500-2000-1500-1000-500-10	elastic distortions	None	0.35 <sup>[18]</sup>	
Convolutional neural network	Committee of 35 conv. net, 1-20-P-40-P-150-10	elastic distortions	Width normalizations	0.23 <sup>[8]</sup>	



### Tensorflow and GPU Acceleration

- Tensorflow, which is developed by Google, can be used to training and inference of deep neural networks.
- □ If you are running the demo on your own computer, you need to install Tensorflow, which can be done in the terminal using command "conda create -n tf tensorflow"
  - It creates a new environment in Anaconda called "tf"
  - · Many standard packages are not in, you need to install scikit-learning, matplotlib etc yourself
- ☐ If you want to use GPU to accelerate Tensorflow on your computer
  - Windows computer with Nvidia graphics card: It should be straightforward
  - MAC with Apple M1/M1 Pro/M1 Max Chip: You can use the build-in Apple GPU
    - Currently, Anaconda does not support Apple ARM based M1 processor
    - Use Miniforge to install: <a href="https://developer.apple.com/metal/tensorflow-plugin/">https://developer.apple.com/metal/tensorflow-plugin/</a> and <a href="https://makeoptim.com/en/deep-learning/tensorflow-plugin/">https://makeoptim.com/en/deep-learning/tensorflow-plugin/</a> and <a href="https://makeoptim.com/en/deep-learning/tensorflow-plugin/">https://makeoptim.com/en/deep-learning/tensorflow-plugin/</a>
- ☐ Or, just use Google Colab
  - You might want to consider to pay for Google Colab Pro and pay a monthly fee (about \$10)
    - After open your .ipynb, choose "Runtime"->"Change runtime type" and choose "Hardware accelerator=GPU/TPU" and "Runtime shape=High-RAM"





# **Downloading MNIST**

```
■MNIST data is available in many sources
import tensorflow as tf

    Note: It has been removed from sklearn

(Xtr,ytr),(Xts,yts) = tf.keras.datasets.mnist.load_data()
                                                          ☐ Tensorflow version:
print('Xtr shape: %s' % str(Xtr.shape))
print('Xts shape: %s' % str(Xts.shape))

    60000 training samples

    10000 test samples

ntr = Xtr.shape[0]
nts = Xts.shape[0]
                                                          ☐ Each sample is a 28 x 28 images
nrow = Xtr.shape[1]
ncol = Xtr.shape[2]
                                                          □ Grayscale: Pixel values \in \{0,1,...,255\}
                                                            ∘ 0 = Black and
Xtr shape: (60000, 28, 28)
                                                            • 255 = White
Xts shape: (10000, 28, 28)
```

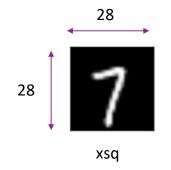


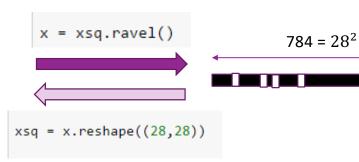
# Matrix and Vector Representation

- $\square$  For this demo, we reshape data from  $N \times 28 \times 28$  to  $N \times 784$
- ☐But, you can easily go back and forth
- □Also, scale the pixel values from -1 to 1

npix = nrow\*ncol
Xtr = 2\*(Xtr/255 - 0.5)
Xtr = Xtr.reshape((ntr,npix))

Xts = 2\*(Xts/255 - 0.5)
Xts = Xts.reshape((nts,npix))





$$S = Mat(x) = \begin{bmatrix} s_{11} & \cdots & s_{1,28} \\ \vdots & \vdots & \vdots \\ s_{28,1} & \cdots & s_{28,28} \end{bmatrix}$$

$$x = \text{vec}(S) = \begin{bmatrix} x_1 & \cdots & x_{784} \end{bmatrix}$$

# Displaying Images in Python









4 random images in the dataset

A human can classify these easily

```
def plt_digit(x):
    nrow = 28
    ncol = 28
    xsq = x.reshape((nrow,ncol))
                                                Key command
    plt.imshow(xsq, cmap='Greys r') ←
    plt.xticks([])
    plt.yticks([])
# Convert data to a matrix
X = mnist.data
y = mnist.target
# Select random digits
                                                Sample
nplt = 4
nsamp = X.shape[0]
                                                permutation is
Iperm = np.random.permutation(nsamp) 
                                                necessary for this
# Plot the images using the subplot command
                                                dataset, as the
for i in range(nplt):
    ind = Iperm[i]
                                                original data is
    plt.subplot(1,nplt,i+1)
                                                ordered by digits
    plt_digit(X[ind,:])
```



# Try a Logistic Classifier

```
ntr1 = 5000
Xtr1 = Xtr[Iperm[:ntr1],:]
ytr1 = ytr[Iperm[:ntr1]]
```

- ☐ Train on 5000 samples
  - To reduce training time.
  - In practice want to train with ~40k
- □ Select correct solver (lbfgs)
  - Others can be very slow. Even this will take minutes

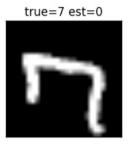


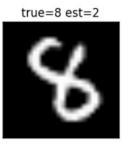
#### Performance

- □Accuracy = 89%. Very bad
- □Some of the errors seem like they should have been easy to spot
- ■What went wrong?

```
nts1 = 5000
Iperm_ts = np.random.permutation(nts)
Xts1 = Xts[Iperm_ts[:nts1],:]
yts1 = yts[Iperm_ts[:nts1]]
yhat = logreg.predict(Xts1)
acc = np.mean(yhat == yts1)
print('Accuaracy = {0:f}'.format(acc))
```

Accuaracy = 0.891000



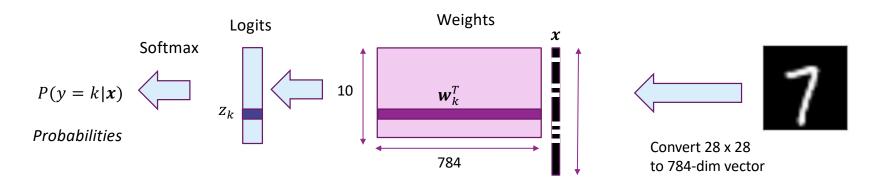








# Recap: Logistic Classifier

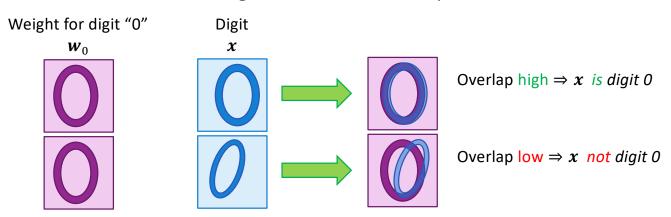


- □ Each logit  $z_k = \mathbf{w}_k^T \mathbf{x}$  = inner product with weight  $\mathbf{w}_k$  with digit  $\mathbf{x}$ , k = 0, ..., 9
- $\square \text{Will select } \hat{y} = \arg \max_{k} P(y = k | x) = \arg \max_{k} z_{k}$ 
  - $\circ$  Output  $z_k$  which is largest
- $\square$  When is  $z_k$  large?



## Interpreting the Logistic Classifier Weights

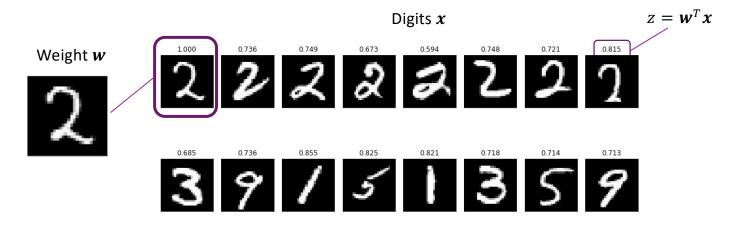
- $\square$  A logit  $z_k = w_k^T x$  is high when there is high overlap between  $w_k$  with digit x
  - Visualize each weight as an image
  - Suppose pixels are 0 or 1
  - $z_k = w_k^T x = \sum_i w_{ki} x_i = \text{number of pixels that overlap with } w_k \text{ and } x_i = w_k^T x_i$
- □ Conclusion: Small variations in digits can cause low overlap





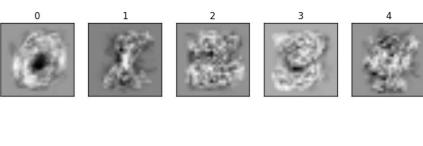
# **Example with Actual Digits**

- ☐ Take weight w from a random digit "2"
- $\square$ Inner products  $z = w^T x$  are only slightly higher for other digits "2"
- $\Box$  Cannot tell which digit is correct from the inner product  $z = w^T x$



# Visualizing the Weights

- □Optimized weights of the classifier
- ☐ Blurry versions of image to try to capture rotations, translations, ...













# Problems with Logistic Classifier

- Linear weighting cannot capture many deformities in image
  - Rotations
  - Translations
  - Variations in relative size of digit components
- ☐ Can be improved with preprocessing
  - E.g. deskewing, contrast normalization, many methods
- ☐ Is there a better classifier?





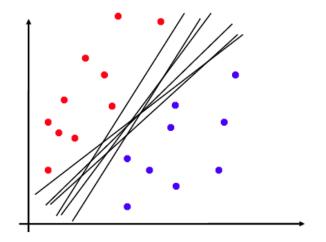
## Outline

- ☐ Motivating example: Recognizing handwritten digits
  - Why logistic regression doesn't work well.
- Maximum margin classifiers
  - ■Support vector machines
  - ☐ Kernel trick
  - lueConstrained optimization



# Non-Uniqueness of Separating Plane

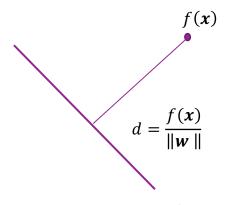
- ☐ Linearly separable data:
  - Can find a separating hyper-plane as a linear classifier.
- ■Separating hyper-plane is not unique
  - Fig. on right: Many separating planes
- ■Which one is optimal?





# Hyperplane Basics

- □Linear function:  $f(x) = w^T x + b, x \in \mathbb{R}^d$
- $\square$  Hyperplane in d-dimensional: f(x) = 0
- □Parameters:
  - Weight w and bias b
  - Unique up to scaling:
  - $\circ$  (b, w) and  $(\alpha b, \alpha w)$  define the same plane.
  - For unique definition, we can require ||w||=1.
- □ Distance of any point **x** to the hyperplane:
  - d = f(x)/||w||, where  $f(x) = b + w^T x$ .
  - ∘ See ESL Sec. 4.5.
  - ESL: Hastie, Tibshirani, Friedman, "The Elements of Statistical Learning". 2<sup>nd</sup> Ed. Springer.



Hyperplane

$$f(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + b = 0$$

# Linear Separability and Margin

- ☐ Given training data  $(x_i, y_i)$ , i = 1, ..., N
  - Binary class label:  $y_i = \pm 1$
- $\square$ Suppose it is separable with parameters (w, b)
- □ There must exist a  $\gamma > 0$  s.t.:

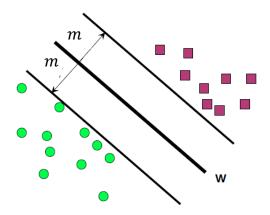
• 
$$b + w_1 x_{i1} + \cdots w_d x_{id} > \gamma$$
 when  $y_i = 1$ 

$$b + w_1 x_{i1} + \cdots w_d x_{id} < -\gamma$$
 when  $y_i = -1$ 

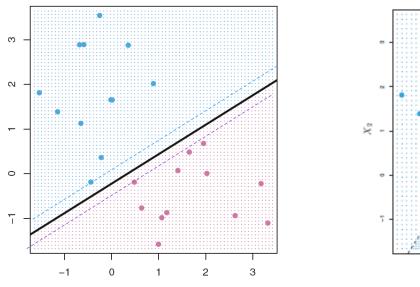
☐ Single equation form:

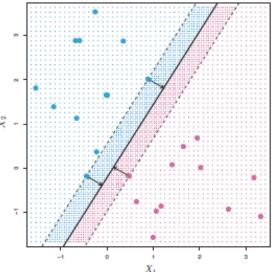
$$y_i(b + w_1x_{i1} + \cdots w_dx_{id}) > \gamma$$
 for all  $i = 1, ..., N$ 

- $\square$  Margin:  $m = \frac{\gamma}{\|w\|}$ : minimal distance of a sample to the plane
  - $\circ$   $\gamma$  is the maximum value satisfying the above constraints



# Which separating plane is better?





From Fig. 9.2 and Fig. 9.3 in ISL.



# Maximum Margin Classifier

- ☐ For the classifier to be more robust to noise, we want to maximize the margin!
- □ Define maximum margin classifier

 $\max_{w,b} \gamma$ 

• Such that  $y_i(b + \mathbf{w}^T \mathbf{x}) \ge \gamma$  for all i Ensures all points are correctly classified

 $\sum_{j=1}^{d} w_j^2 \le 1 \quad \blacksquare$ 

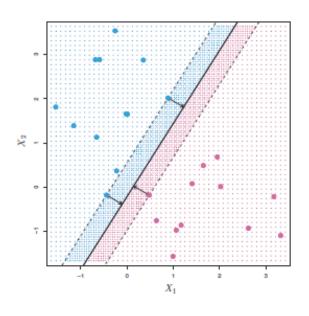
Maximizes the margin

Scaling on weights

- □ Called a constrained optimization
  - Objective function and constraints
  - More on this later.
- See closed form solution in Sec. 4.5.2 in ESL. Note notation difference.



# Visualizing Maximum Margin Classifier



- ☐Fig. 9.3 of ISL
- ☐ Margin determined by closest points to the line
  - The maximal margin hyperplane represents the midline of the widest "slab" that we can insert between two classes
- ☐ In this figure, there are 3 points at the margin

ISL: James, Witten, Hastie, Tibshirani, An Introduction to Statistical Learning, Springer. 2013.

## Problems with MM classifier

- ☐ Data is often not perfectly separable
  - Only want to correctly separate most points

- ☐MM classifier is not robust
  - A single sample can radically change line

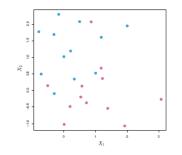
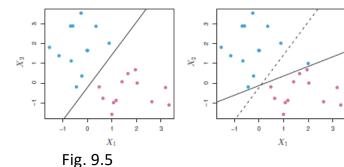


Fig. 9.4



#### **In-Class Exercise**

☐ Found in github site: svm\_inclass.ipynb

#### Problem 1. Margin

For the points below with binary labels:

- . Create a scatter plot of the points with different markers for the two classes
- . Find the weight and bias of the classifier that separates the two classes
- . Compute the distance to the classifier boundary for the points
- · Find the margin

```
X = np.array([[0.5,0.5], [1,0.5],[0.5,1.75], [0.75,2.75], [1.1,2.2], [2,1], [3,1.5]])
y = np.array([1,1,1,0,0,0,0])
```





## Outline

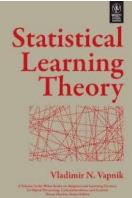
- ☐ Motivating example: Recognizing handwritten digits
  - Why logistic regression doesn't work well.
- ☐ Maximum margin classifiers
- Support vector machines
- ☐ Kernel trick
- lueConstrained optimization



# Support Vector Machine

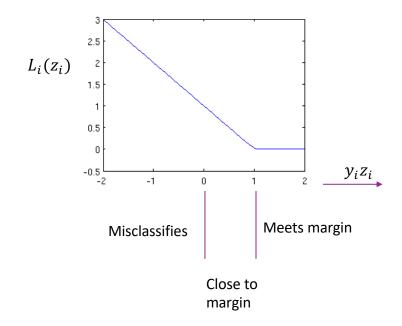
- ■Support Vector Machine (SVM)
  - Vladimir Vapnik, 1963
  - But became widely-used with kernel trick, 1993
  - More on this later
- ☐Got best results on character recognition
- ☐ Key idea: Allow "slack" in the classification
  - Support vector classifier (SVC): Directly use raw features.
     Good when the original feature space is roughly linearly separable
  - Support vector machine (SVM): Map the raw features to some other domain through a kernel function





# Hinge Loss

- $\square$  Fix  $\gamma = 1$
- □Want ideally:  $y_i(\mathbf{w}^T\mathbf{x} + b) \ge 1$  for all samples i
  - Equivalently,  $y_i z_i \ge 1$ ,  $z_i = b + \mathbf{w}^T \mathbf{x}$
- ☐But perfect separation may not be possible
- ☐ Define hinge loss or soft margin:
  - $L_i(\mathbf{w}, b) = \max(0, 1 y_i z_i)$
- ☐ Starts to increase as sample is misclassified:
  - $y_i z_i \geq 1 \ \Rightarrow \ {
    m Sample meets margin target}, \ L_i(w) = 0$
  - ∘  $y_i z_i \in [0,1)$  ⇒ Sample margin too small, small loss
  - $y_i z_i \le 0 \Rightarrow$  Sample misclassified, large loss



# **SVM Optimization**

- $\square$  Given data  $(x_i, y_i)$

$$J(\boldsymbol{w}, b) = C \sum_{i=1}^{N} \max(0, 1 - y_i(\boldsymbol{w}^T \boldsymbol{x}_i + b)) + \frac{1}{2} \|\boldsymbol{w}\|^2$$
rols final margin. Hinge loss term. margin=1/\|\boldsymbol{w}\|

C controls final margin

Hinge loss term Attempts to reduce

Misclassifications

- $\square$ Constant C > 0 will be discussed below
- □Note: ISL book uses different naming conventions.
  - We have followed convention in sklearn



# Alternate Form of SVM Optimization

☐ Equivalent optimization:

$$\min J_1(\boldsymbol{w}, b, \boldsymbol{\epsilon}), \qquad J_1(\boldsymbol{w}, b, \boldsymbol{\epsilon}) = C \sum_{i=1}^N \epsilon_i + \frac{1}{2} \|\boldsymbol{w}\|^2$$

■Subject to constraints:

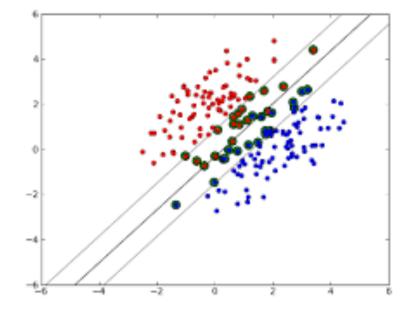
$$y_i(\mathbf{w}^T \mathbf{x}_i + b) \ge 1 - \epsilon_i \text{ for all } i = 1, ..., N$$
  
 $\epsilon_i \ge 0 \text{ for all } i = 1, ..., N$ 

- $\epsilon_i$  = amount sample *i* misses margin target
- $\square$  Sometimes write as  $J_1(w, b, \epsilon) = C \|\epsilon\|_1 + \frac{1}{2} \|w\|^2$ 
  - $\|\epsilon\|_1 = \sum_{i=1}^N \epsilon_i$  called the "one-norm"
  - $\circ\,$  Generally one-norm would have absolute sign over  $\varepsilon_i.$
  - But in this case, when the constraint is met,  $\epsilon_i >= 0$ .

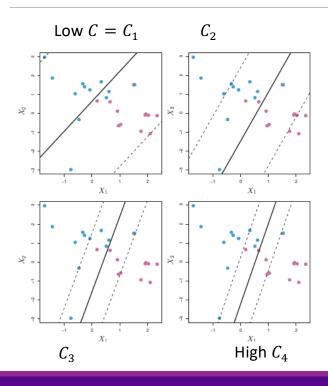


# **Support Vectors**

- □Support vectors: Samples that either:
  - Are exactly on margin:  $y_i(\mathbf{w}^T \mathbf{x}_i + b) = 1$
  - Or, on wrong side of margin:  $y_i(\mathbf{w}^T \mathbf{x}_i + b) \leq 1$
- ☐ Changing samples that are not SVs
  - Does not change solution
  - Provides robustness



# Illustrating Effect of C



#### ☐Fig. 9.7 of ISL

- $\circ$  Note: C has opposite meaning in ISL than python
- Here, we use python meaning

#### $\square$ Low C:

- Leads to large margin
- But allow many violations of margin.
- Many more SVs
- Reduces variance by using more samples

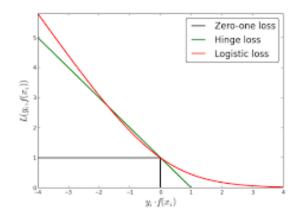
#### ☐ Large C:

- Leads to small margin
- Reduce number of violations, and fewer SVs.
- Highly fit to data. Low bias, higher variance
- More chance to overfit

# Relation to Logistic Regression

□ Logistic regression also minimizes a loss function:

$$J(\mathbf{w}, b) = \sum_{i=1}^{N} L_i(\mathbf{w}, b), \qquad L_i(\mathbf{w}, b) = \ln P(y_i | \mathbf{x}_i) = -\ln(1 + e^{-y_i z_i})$$



#### **In-Class Exercise**

#### **Problem 2. Minimizing the Hinge Loss**

For the data below, first create a scatter plot of the points with different markers for the two classes. You should see that the data is not linearly separable.

Then, consider a set of classifiers:

```
yhat = sign(z), z = w.dot(x)+b
```

Use the the w below, plot the hinge loss as a function of the bias b where the hinge loss is:

```
J = sum( maximum(0, 1-ypm*z) )
```

Here ypm=2\*y-1 so that it is a value +1 or -1. Find the b that minimizes the hinge loss and plot the boundary of the classifier.

```
X = np.array([[0.5,0.5], [1,0.5],[0.5,1.75], [2,2], [0.75,0.75], [0.75,2.75], [1.1,2.2], [2,1], [3,1.5]])
y = np.array([1,1,1,1,0,0,0,0,0])

w = np.array([1.5, 1])
w = w / np.linalg.norm(w)
```





## Outline

- ☐ Motivating example: Recognizing handwritten digits
  - Why logistic regression doesn't work well.
- ☐ Maximum margin classifiers
- ■Support vector machines
- Kernel trick
- ☐ Constrained optimization



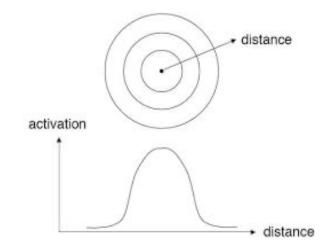
## The Kernel Function

#### □Kernel function:

- Function  $K(x_i, x)$
- Key function for SVMs and kernel classifiers
- $\circ$  Measures "similarity" between new sample x and training sample  $x_i$

#### ☐ Typical property

- $x_i, x \text{ close} \Rightarrow K(x_i, x) \text{ maximum value}$
- $x_i, x \text{ far} \Rightarrow K(x_i, x) \approx 0$



### Common Kernels

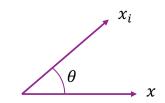
#### ☐Linear SVM:

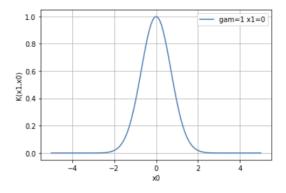
- $K(x_i, x) = x_i^T x = ||x_i|| ||x|| \cos \theta$
- Maximum when angle between vectors is small
- □ Radial basis function:

$$K(x_i, x) = \exp[-\gamma ||x - x_i||^2]$$

 $\circ$  1/ $\gamma$  indicates width of kernel

• Typically d=2

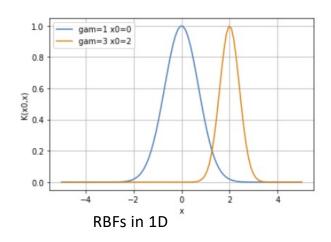


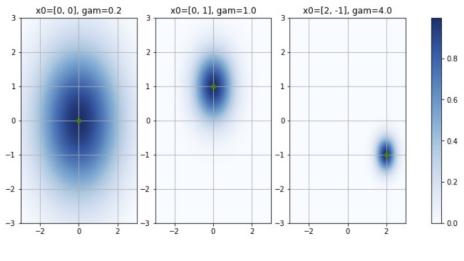


# **RBF Kernel Examples**

 $\square$ RBF kernel:  $K(x_0, x) = \exp[-\gamma ||x - x_0||^2]$ 

- $\circ$  Peak value of 1 at  $x = x_0$
- Width  $\propto \frac{1}{\gamma}$







### Kernel Classifier

#### **□**Given:

- Training data  $(x_i, y_i)$  with binary labels  $y_i = \pm 1$
- Kernel  $K(x_i, x)$

#### $\square$ To classify a new point x:

- Decision function:  $z = \sum_{i=1}^{n} y_i K(x_i, x)$
- Classify:  $\hat{y} = sign(z)$

#### ☐Idea:

- $\circ z$  is large positive when x is close to samples  $x_i$  with  $y_i = 1$
- $\circ z$  is large negative when x is close to samples  $x_i$  with  $y_i = -1$
- ☐ Kernel classifiers are a subject on their own
  - We just mention them here to explain connection to SVMs

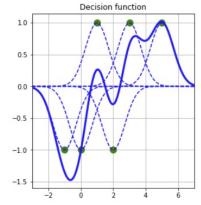


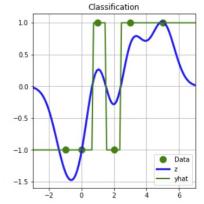
# Example in 1D

- $\square$  Example data with 6 points  $(x_i, y_i)$ 
  - RBF kernel:  $K(x_i, x) = e^{-\gamma(x_i x)^2}$ ,  $\gamma = 1$
- □ Decision function:
  - $\circ \ z = \sum_{i=1}^n y_i K(x_i, x)$
  - Sum of bell curves
  - Positive when near positive samples
  - Negative when near negative samples
- □Classification:

$$\hat{y} = sign(z)$$

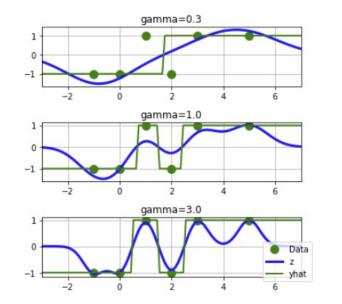
i	1	2	3	4	5	6
$x_i$	-1	0	1	2	3	5
$y_i$	-1	-1	1	-1	1	1





### **Effect of Gamma**

- ■Same data as before
- $\square$ RBF kernel:  $K(x_i, x) = e^{-\gamma(x_i x)^2}$
- $\square$  As  $\gamma$  increases:
  - Decision function  $z \approx y_i$  when  $x = x_i$
  - Classifier fits training data better
  - Classification region more complex
- $\square$  As a classifier, higher  $\gamma$  results in:
  - Lower bias error
  - But, higher variance error



### **SVMs with Non-Linear Transformations**

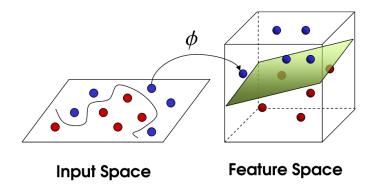
- Non-linear transformation:
  - Replace x with  $\phi(x)$
  - Enables more rich, non-linear classifiers
  - Examples: polynomial classification

$$\phi(x) = [1, x, x^2, \dots, x^{d-1}]$$

☐ Tries to find separation in a feature space



Makes applying non-linear transformations easy



### SVM with the Transformation

- $\square$  Consider SVM model with x replaced by  $\phi(x)$
- ☐ Minimize SVM cost function as before (i.e. Hinge loss + inverse margin)
- ☐ Theorem: The optimal weight is of the form:

$$w = \sum_{i=1}^{N} \alpha_i y_i \phi(x_i)$$

- $\alpha_i \geq 0$  for all i
- $\alpha_i > 0$  if and only if sample *i* is a support vector
- Will show this fact later using results in constrained optimization
- $\square$  Consequence: The linear discriminant on any other sample x is:

$$z = b + \mathbf{w}^{T} \phi(\mathbf{x}) = b + \sum_{i=1}^{N} \alpha_{i} y_{i} \boxed{\phi(\mathbf{x}_{i})^{T} \phi(\mathbf{x})} \qquad K(\mathbf{x}_{i}, \mathbf{x}) = \text{"kernel"}$$

0





### Kernel Form of the SVM Classifier

□SVM classifier can be written with the kernel  $K(x_i, x)$  and values  $\alpha_i \ge 0$ :

$$z = b + \sum_{i=1}^{N} \alpha_i y_i K(x_i, x),$$

$$\hat{y} = \text{sign}(z) = \begin{cases} 1 & \text{if } z > 0 \\ -1 & \text{if } z < 0 \end{cases}$$
Classification decision

- □ Key point: SVM classifier is approximately Kernel classifier
- ■But there are two differences:
  - Weights  $\alpha_i \geq 0$  on the samples (the weights are only non-zero on the SVs)
  - A bias term *b* (can be positive or negative)



### "Kernel Trick" and Dual Parameterization

□Kernel form of SVM classifier (previous slide):

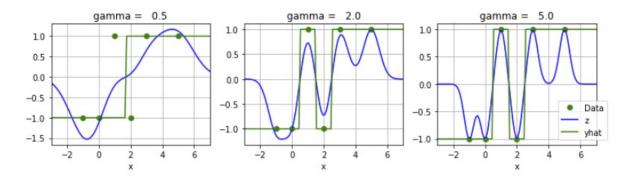
$$z = b + \sum_{i=1}^{N} \alpha_i y_i K(\mathbf{x}_i, \mathbf{x}),$$
  
$$\hat{y} = \operatorname{sign}(z)$$

- □ Dual parameters:  $\alpha_i \ge 0$ , i = 1, ..., N
  - Called the dual parameters due to constrained optimization see next section
- ☐Kernel trick:
  - $\circ$  Directly solve the parameters  $\alpha$  instead of the weights w
  - $\circ$  Can show that the optimization only needs the kernel  $K(x_i, x)$
  - $\circ$  Does not need to explicitly use  $\phi(x)$

# SVM Example in 1D

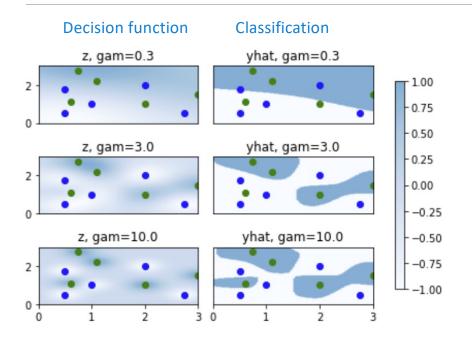
- ☐Same data as in the Kernel classifier example
- $\square$  Fit SVM with RBF with different  $\gamma$
- $\square$ Similar trends as kernel classifier: As  $\gamma$  increases
  - z "fits" data  $(x_i, y_i)$  closer
  - Leads to more complex decision regions.
  - Enables nonlinear decision regions

i	1	2	3	4	5	6
$x_i$	-1	0	1	2	3	5
$y_i$	-1	-1	1	-1	1	1





# Example in 2D



#### **□**Example:

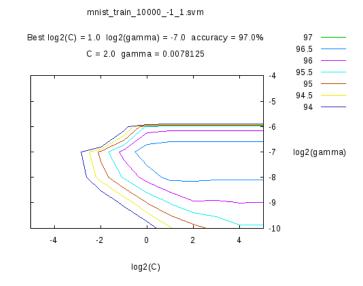
- 10 data points with binary labels
- $\circ$  Fit SVM with C=1 and RBF
- $\circ$   $\gamma = 0.3$ , 3 and 10

#### ■Plot:

- ∘ *z*= linear discriminant
- $\hat{y} = sign(z) = classification decision$
- $\square$  Observe: As  $\gamma$  increases
  - Fits training data better
  - More complex decision region

#### Parameter Selection

- ☐ For SVMs with RBFs we need to select:
  - $\circ$  Parameter C > 0 in the loss function
  - $\circ$  Kernel width  $\gamma > 0$
- $\square$  Higher C or  $\gamma$ 
  - Fewer SVs
  - Classifiers averages over smaller set
  - Lower bias, but higher variance
- ☐ Typically select via cross-validation
  - $\circ$  Try out different  $(C, \gamma)$  pairs
  - Find which one provides highest accuracy on test set
- ☐ Python can automatically do grid search



http://peekaboo-vision.blogspot.com/2010/09/mnist-for-ever.html





### Multi-Class SVMs

- $\square$ Suppose there are K classes
- One-vs-one:
  - Train  $\binom{K}{2}$  SVMs for each pair of classes
  - Test sample assigned to class that wins "majority of votes"
  - Best results but very slow
- □One-vs-rest:
  - $\circ$  Train K SVMs: train each class k against all other classes
  - $\,^{\circ}\,$  Pick class with highest  $z_k$
- ☐Sklearn has both options



#### **MNIST Results**

- □ Run classifier
- ■Very slow
  - Several minutes for 40,000 samples
  - Slow in training and test
  - Major drawback of SVM
- $\square$ Accuracy  $\approx 0.984$ 
  - Much better than logistic regression
- □Can get better with:
  - pre-processing
  - More training data
  - Optimal parameter selection

```
from sklearn import svm

# Create a classifier: a support vector classifier
svc = svm.SVC(probability=False, kernel="rbf", C=2.8, gamma=.0073,verbose=10)

svc.fit(Xtr,ytr)

[LibSVM]

SVC(C=2.8, cache_size=200, class_weight=None, coef0=0.0,
    decision_function_shape=None, degree=3, gamma=0.0073, kernel='rbf',
    max_iter=-1, probability=False, random_state=None, shrinking=True,
    tol=0.001, verbose=10)

yhat1 = svc.predict(Xts)
acc = np.mean(yhat1 == yts)
print('Accuaracy = {0:f}'.format(acc))
```

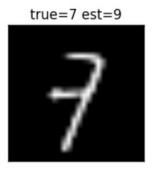
Accuaracy = 0.984000

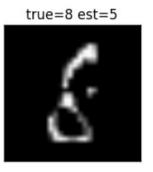


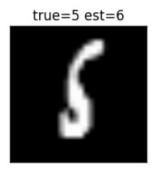


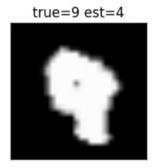
## **MNIST Errors**

■Some of the error are hard even for a human









## Outline

- ☐ Motivating example: Recognizing handwritten digits
  - Why logistic regression doesn't work well.
- ☐ Maximum margin classifiers
- ■Support vector machines
- ☐ Kernel trick
- Constrained optimization



# **Constrained Optimization**

- ☐ In many problems, variables are constrained
- □ Constrained optimization formulation:
  - Objective: Minimize f(w)
  - Constraints:  $g_1(\mathbf{w}) \leq 0, ..., g_M(\mathbf{w}) \leq 0$
- **■**Examples:
  - Minimize the mpg of a car subject to a cost or meeting some performance
  - In ML: weight vector may have constraints from physical knowledge
- $\square$  Often write constraints in vector form: Write  $g(\mathbf{w}) \leq 0$

$$g(\mathbf{w}) = [g_1(\mathbf{w}), \dots, g_m(\mathbf{w})]^T$$



## Lagrangian

- □ Constrained optimization: Min f(w) s.t.  $g(w) \le 0$
- $\square$  Consider first a single constraint: g(w) is a scalar
- $\Box \text{Define Lagrangian: } L(\mathbf{w}, \lambda) = f(\mathbf{w}) + \lambda g(\mathbf{w})$ 
  - w is called the primal variable
  - $\circ$   $\lambda$  is called the dual variable
- $\square$  Dual minimization: Given a dual parameter  $\lambda$ , minimize

$$\widehat{\boldsymbol{w}}(\lambda) = \arg\min_{\boldsymbol{w}} L(\boldsymbol{w}, \lambda), \qquad L^*(\lambda) = \min_{\boldsymbol{w}} L(\boldsymbol{w}, \lambda)$$

- Minimizes a weighted combination of objective and constraint.
- Higher  $\lambda \Rightarrow$  Weight constraint more (try to make  $g(\mathbf{w})$  smaller)
- Lower  $\lambda \Rightarrow$  Weight objective more (try to make f(w) smaller)



### **KKT Conditions**

- $\square$  Given objective f(w) and constraint g(w)
- $\square$ KKT Conditions:  $\widehat{\boldsymbol{w}}$ ,  $\widehat{\lambda}$  satisfy:
  - $\widehat{\boldsymbol{w}}$  minimizes the Lagrangian:  $\widehat{\boldsymbol{w}} = \arg\min_{\boldsymbol{w}} L(\boldsymbol{w}, \widehat{\lambda})$
  - Either
    - $g(\widehat{\mathbf{w}}) = 0$  and  $\widehat{\lambda} \ge 0$  [active constraint]
    - $g(\hat{w}) < 0$  and  $\hat{\lambda} = 0$  [inactive constraint]
- ☐ Theorem: Under some technical conditions,
  - $\circ$  if  $\hat{\boldsymbol{w}}$ ,  $\hat{\lambda}$  are local mimima of the constrained optimization, they must satisfy KKT conditions



# General Procedure for Single Constraint

#### ■Suppose:

- $\mathbf{w} = (w_1, ..., w_d)^T$ : d unknown primal variables
- $g(\mathbf{w}) \leq 0$ : scalar constraint

#### □ Case 1: Assume constraint is active:

- Solve w and  $\lambda$ :  $\partial L(w,\lambda)/\partial w_i=0$  and g(w)=0 (resulting from setting  $\partial L(w,\lambda)/\partial \lambda=0$ )
- d+1 unknowns and d+1 equations
- $\,^{\circ}\,$  Verify that  $\lambda \geq 0$

#### □ Case 2: Assume constraint is inactive

- Solve primal objective  $\partial f(\mathbf{w})/\partial w_i = 0$  ignoring constraint
- $\circ d$  unknowns and d equations
- Verify that constraint is satisfied:  $g(\mathbf{w}) \leq 0$



### KKT Conditions Illustrated

■Example 1: Constraint is "active"

$$\min_{w} w^2 \quad s. t. \ w + 1 \le 0$$

☐ Example 2: Constraint is "inactive"

$$\min_{w} w^2 \quad s.t. \ w - 1 \le 0$$

☐ Examples worked on board with illustration



## Multiple Constraints

- □ Now consider constraint:  $g(\mathbf{w}) = [g_1(\mathbf{w}), ..., g_M(\mathbf{w})]^T \le 0$ .
- □ Lagrangian is:

$$L(\mathbf{w}, \lambda) = f(\mathbf{w}) + \lambda^{T} g(\mathbf{w}) = f(\mathbf{w}) + \sum_{m=1}^{M} \lambda_{m} g_{m}(\mathbf{w})$$

- Weighted sum of all *M* constraints
- $\circ$   $\lambda$  is called the dual vector
- □KKT conditions extend to:
  - $\hat{w}$  minimizes the Lagrangian:  $\hat{w} = \arg\min_{w} L(w, \hat{\lambda})$
  - $\circ$  For each m = 1, ..., M
    - $g_m(\widehat{\boldsymbol{w}}) = 0$  and  $\widehat{\lambda}_m \ge 0$  [active constraint]
    - $g_m(\widehat{\boldsymbol{w}}) < 0$  and  $\hat{\lambda}_m = 0$  [inactive constraint]



# Multiple Constraints

- $\square$  If there are M constraints, there could be  $2^M$  cases to discuss
- ☐ In practice, the number of cases are much smaller
- ☐ For more information on KKT conditions, check the following lecture on youtube
  - UAMathCamp Lecture 40(A): Kuhn-Tucker Conditions: Conceptual and geometric insight
  - https://www.youtube.com/watch?v=HIm3Z0L90Co



## **SVM Constrained Optimization**

☐ Recall: SVM constrained optimization

$$\min J_1(\boldsymbol{w}, b, \boldsymbol{\epsilon}), \qquad J_1(\boldsymbol{w}, b, \boldsymbol{\epsilon}) = C \sum_{i=1}^N \epsilon_i + \frac{1}{2} \|\boldsymbol{w}\|^2$$

- Constraints:  $y_i(\mathbf{w}^T \mathbf{x}_i + b) \ge 1 \epsilon_i$  and  $\epsilon_i \ge 0$  for all i = 1, ..., N
- ☐ After applying KKT conditions and some algebra [beyond this class], solution is
  - Optimal weight vector:  $\mathbf{w} = \sum_{i=1}^{N} \alpha_i y_i \mathbf{x}_i$  linear combination of instances
  - Dual parameters  $\alpha_i$  minimize

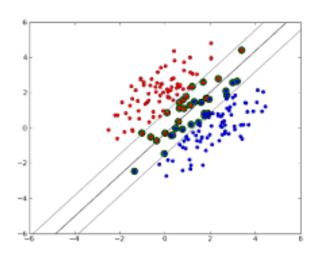
$$\sum_{i=1}^{N} \alpha_i - \frac{1}{2} \sum_{i=1}^{N} \sum_{j=1}^{N} \alpha_i \alpha_j y_i y_j \mathbf{x}_i^T \mathbf{x}_j \quad \text{s.t. } 0 \le \alpha_i \le C$$

- Details can be found on textbook ESLII, section 12.2.1 Computing the Support Vector Classifier
- MIT 6.034 Artificial Intelligence, Fall 2010 https://www.youtube.com/watch?v=\_PwhiWxHK8o



# **Support Vectors**

- lacksquare Classifier weight is:  $m{w} = \sum_{i=1}^N \alpha_i y_i m{x}_i$
- lacksquare Can show that  $lpha_i>0$  only when  $x_i$  is a support vector
  - On boundary or violating constraint
  - $\circ$  Otherwise  $\alpha_i = 0$





## What you should know

- □ Interpret weights in linear classification of images (logistic regression): Match filters
- ☐ Understand the margin in linear classification and maximum margin classifier
- □SVM classifier: Allow violation of margin by introducing slack variables (More robust than linear classifier)
- □ Solve constrained optimization using the Lagrangian.
  - Understand KKT conditions for a single constraint
- Extend to nonlinear classifier by feature transformation: SVM with nonlinear kernels
- □ Select SVM parameters from cross-validation



