

Introduction to Machine Learning

Unit 4 Solutions: Model Order Selection

Prof. Sundeep Rangan

1. For each of the following pairs of true functions $f_0(\mathbf{x})$ and model classes $f(\mathbf{x}, \boldsymbol{\beta})$ determine: (i) if the model class is linear; (ii) if there is no under-modeling; and (iii) if there is no under-modeling, what is the true parameter?

(a) $f_0(x) = 1 + 2x$, $f(x, \boldsymbol{\beta}) = \beta_0 + \beta_1 x + \beta_2 x^2$

(b) $f_0(x) = 1 + 1/(2 + 3x)$, $f(x, a_0, a_1, b_0, b_1) = (a_0 + a_1 x)/(b_0 + b_1 x)$.

(c) $f_0(x) = (x_1 - x_2)^2$ and

$$f(\mathbf{x}, a, b_1, b_2, c_1, c_2) = a + b_1 x_1 + b_2 x_2 + c_1 x_1^2 + c_2 x_2^2.$$

Solution

(a) Linear model; no undermodeling; true parameter is $\boldsymbol{\beta} = (1, 2, 0)$.

(b) The model is nonlinear. We can write

$$f_0(x) = 1 + \frac{1}{2 + 3x} = \frac{3 + 3x}{2 + 3x},$$

So, there is no under-modeling and the true parameters are $(a_0, a_1, b_0, b_1) = (3, 3, 2, 3)$.

(c) The model is linear. The true function is

$$f_0(x) = (x_1 - x_2)^2 = x_1^2 - 2x_1 x_2 + x_2^2.$$

There is undermodeling since the model class doesn't have an $x_1 x_2$ term.

2. You want to fit an exponential model of the form,

$$y \approx \hat{y} = \sum_{j=0}^d \beta_j e^{-ju/d},$$

where the input u and output y are scalars. You are given python functions:

```
model = LinearRegression()
model.fit(X,y)           # Fits a linear model for a data matrix X
yhat = model.predict(X)  # Predicts values
```

Using these functions, write python code that, given vectors \mathbf{u} and \mathbf{y} :

- Splits the data into training and test using half the samples for each.
- Fits models of order $d_{\text{test}} = [1, 2, \dots, 10]$ on the training data.
- Selects the model with the lowest mean squared error.

Solution One solution is as follows:

```
# Split data into training and test
# You can also use train_test_split from sklearn
n = len(u)
ntr = n // 2
utr = u[:ntr]
ytr = y[:ntr]
uts = u[ntr:]
yts = y[ntr:]

# Loop over model orders
dtest = np.arange(1,11) # Note this loop goes d=1,...,10
nd = len(dtest)
mse = np.zeros(nd)
for i, d in dtest:

    # Transform the data
    # Create the feature matrices using python broadcasting
    # Xtr[i,j] = np.exp(-utr[i]*j/d)
    # Xts[i,j] = np.exp(-uts[i]*j/d)
    powers = np.arange(0,d+1)/d
    Xtr = np.exp(-utr[:,None]*powers[None,:])
    Xts = np.exp(-uts[:,None]*powers[None,:])

    # Create model
    model = LinearRegression()

    # Fit model on training data
    model.fit(Xtr, ytr)

    # Measure MSE on test data
    yhat = model.predict(Xts)
    mse[i] = np.mean((yhat - yts)**2)

# Select model with lower test error
im = np.argmin(mse)
dopt = dtest[im]
```

3. Suppose we want to fit a model,

$$y \approx \hat{y} = f(x, \beta) = \beta x^2.$$

We get data (x_i, y_i) , $i = 1, \dots, N$ and compute the estimate,

$$\hat{\beta} = \frac{\sum_{i=1}^N y_i}{\sum_{i=1}^N x_i^2}.$$

Note: This is not optimal least-squares estimator. But, it is easier to analyze. For each case below compute the bias,

$$\text{Bias}(x) := \mathbb{E}(f(x, \hat{\beta})) - f(x, \beta_0),$$

as a function of the test point x , true parameter β_0 and test data x_i .

- (a) The training data has no noise: $y_i = f(x_i, \beta_0)$.
- (b) The training data is $y_i = f(x_i, \beta_0) + \epsilon_i$ where the noise is i.i.d. $\epsilon_i \sim \mathcal{N}(0, \sigma^2)$.
- (c) The training data is $y_i = f(x_i + \epsilon_i, \beta_0)$ where the noise is i.i.d. $\epsilon_i \sim \mathcal{N}(0, \sigma^2)$.

Solution To compute the bias, first observe that

$$\begin{aligned} \text{Bias}(x) &:= \mathbb{E}(f(x, \hat{\beta})) - f(x, \beta_0) = \mathbb{E}(\hat{\beta}x^2) - \beta_0x^2 \\ &= [\mathbb{E}(\hat{\beta}) - \beta_0]x^2, \end{aligned} \tag{1}$$

where we have used the linearity of expectation. Remember that the test point x is not random, so you can “pull it out” of the expectation. Now, we look at the expectation of $\hat{\beta}$:

$$\mathbb{E}(\hat{\beta}) = \mathbb{E}\left[\frac{\sum_{i=1}^N y_i}{\sum_{i=1}^N x_i^2}\right] = \frac{\sum_{i=1}^N \mathbb{E}(y_i)}{\sum_{i=1}^N x_i^2}, \tag{2}$$

where we have used that training data x_i is not random. So, again we can pull the denominator out of the expectation. We can now evaluate the bias for each of the three cases.

- (a) In this case,

$$\mathbb{E}(y_i) = \mathbb{E}(f(x_i, \beta_0)) = \mathbb{E}(x_i^2 \beta_0) = x_i^2 \beta_0.$$

In the last step, since there is no randomness (remember the training data x_i and true parameter β_0 are not random). Substituting this expectation into (2) we obtain,

$$\mathbb{E}(\hat{\beta}) = \frac{\sum_{i=1}^N \mathbb{E}(y_i)}{\sum_{i=1}^N x_i^2} = \frac{\sum_{i=1}^N x_i^2 \beta_0}{\sum_{i=1}^N x_i^2} = \beta_0.$$

Therefore, from (1),

$$\text{Bias}(x) = [\mathbb{E}(\hat{\beta}) - \beta_0]x^2 = [\beta_0 - \beta_0]x^2 = 0.$$

So, the bias is zero. We say the estimator is *unbiased*.

- (b) In this case,

$$\begin{aligned} \mathbb{E}(y_i) &= \mathbb{E}(f(x_i, \beta_0) + \epsilon_i) = \mathbb{E}(x_i^2 \beta_0 + \epsilon_i) \\ &= x_i^2 \beta_0 + \mathbb{E}(\epsilon_i) = x_i^2 \beta_0, \end{aligned}$$

where in the last step, we used that $\epsilon_i \sim \mathcal{N}(0, \sigma_i^2)$ so $\mathbb{E}(\epsilon_i) = 0$. The final expression, $\mathbb{E}(y_i) = x_i^2 \beta_0$ is identical to part (a). So, again we get $\text{Bias}(x) = 0$ for all test points x .

(c) For this case,

$$\begin{aligned}\mathbb{E}(y_i) &= \mathbb{E}(f(x_i + \epsilon_i, \beta_0)) = \mathbb{E}((x_i + \epsilon_i)^2 \beta_0) \\ &= [x_i^2 \beta_0 + 2\mathbb{E}(\epsilon_i)x_i \beta_0 + \mathbb{E}(\epsilon_i^2)] \beta_0 \\ &= [x_i^2 + \sigma^2] \beta_0,\end{aligned}$$

where, in the last step, we used that $\mathbb{E}(\epsilon_i) = 0$ and $\mathbb{E}(\epsilon_i^2) = \sigma^2$. Substituting this into (2),

$$\mathbb{E}(\hat{\beta}) = \frac{\sum_{i=1}^N \mathbb{E}(y_i)}{\sum_{i=1}^N x_i^2} = \frac{\sum_{i=1}^N (x_i^2 + \sigma^2) \beta_0}{\sum_{i=1}^N x_i^2} = \beta_0 + \frac{\beta_0 N \sigma^2}{\sum_{i=1}^N x_i^2}.$$

Substituting into (1),

$$\text{Bias}(x) = [\mathbb{E}(\hat{\beta}) - \beta_0] x^2 = \frac{\beta_0 N \sigma^2 x^2}{\sum_{i=1}^N x_i^2}.$$

So, in this case, there is a bias in the estimator.

4. In this problem, we will see how to calculate the bias when there is undermodeling. Suppose that training data (x_i, y_i) , $i = 1, \dots, n$ is fit using a simple linear model of the form,

$$\hat{y} = f(x, \beta) = \beta_0 + \beta_1 x.$$

However, the true relation between x and y is given

$$y = f_0(x), \quad f_0(x) = \beta_{00} + \beta_{01}x + \beta_{02}x^2,$$

where the “true” function $f_0(x)$ is quadratic and $\beta_0 = (\beta_{00}, \beta_{01}, \beta_{02})$ is the vector of the true parameters. There is no noise.

- Write an expression for the least-squares estimate $\hat{\beta} = (\hat{\beta}_0, \hat{\beta}_1)$ in terms of the training data (x_i, y_i) , $i = 1, \dots, n$. These expressions will involve multiple steps. You do not need to simplify the equations. Just make sure you state clearly how one would compute $\hat{\beta}$ from the training values.
- Using the fact that $y_i = f_0(x_i)$ in the training data, write the expression for $\beta = (\hat{\beta}_0, \hat{\beta}_1)$ in terms of the values x_i and the true parameter values β_0 . Again, you do not need to simplify the equations. Just make sure you state clearly how one would compute $\hat{\beta}$ from the true parameter vector β_0 and \mathbf{x} .
- Suppose that the true parameters are $\beta_0 = (1, 2, -1)$ and the model is trained using 10 values x_i uniformly spaced in $[0, 1]$. Write a short python program to compute the estimate parameters $\hat{\beta}$. Plot the estimated function $f(x, \hat{\beta})$ and true function $f_0(x)$ for $x \in [0, 3]$.
- For what value x in this range $x \in [0, 3]$ is the bias $\text{Bias}^2(x) = (f(x, \hat{\beta}) - f_0(x))^2$ largest?

Solution

- (a) This is simple linear regression so, from the class notes, the parameter estimates are

$$\hat{\beta}_1 = \frac{s_{xy}}{s_{xx}}, \quad \beta_0 = \bar{y} - \hat{\beta}_1 \bar{x},$$

where \bar{x} and \bar{y} are the sample means and s_{xy} and s_{xx} are the sample co-variance and variance:

$$\bar{x} = \frac{1}{n} \sum_{i=1}^n x_i, \quad \bar{y} = \frac{1}{n} \sum_{i=1}^n y_i,$$

$$s_{xy} = \frac{1}{n} \sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y}), \quad s_{xx} = \frac{1}{n} \sum_{i=1}^n (x_i - \bar{x})^2.$$

- (b) Use the above expressions but substitute $y_i = \beta_{00} + \beta_{01}x_i + \beta_{02}x_i^2$.
- (c) Suppose that the true parameters are $\beta_0 = (1, 2, 0.5)$ and the model is trained using 10 values x_i uniformly spaced in $[0, 1]$. Write a short python program to compute the estimate parameters $\hat{\beta}$. Plot the estimated function $f(x, \hat{\beta})$ and true function $f_0(x)$ for $x \in [0, 3]$.

You can compute the estimate and plot the estimate and true function with the following code:

```
import numpy.polynomial.polynomial as poly

beta0 = np.array([1,2,-1]) # True parameter value
x = np.linspace(0,1,10)    # Training values for x
y = poly.polyval(x,beta0)  # Training values for y

# Get parameter estimate based on simple linear regression formula
xm = np.mean(x)
ym = np.mean(y)
sxy = np.mean((x-xm)*(y-ym))
sxx = np.mean((x-xm)**2)
betahat1 = sxy/sxx
betahat0 = ym - betahat1*xm

# Plot true function and estimate
xp = np.linspace(0,3,100)
yp0 = poly.polyval(xp,beta0)
yphat = betahat0 + betahat1*xp

plt.plot(xp,np.column_stack((yp0, yphat)), '-')
plt.scatter(x,y)
plt.legend(['True', 'Est', 'Training'], loc='upper left')
plt.grid()
plt.xlim([0,3])
plt.xlabel('x')
```

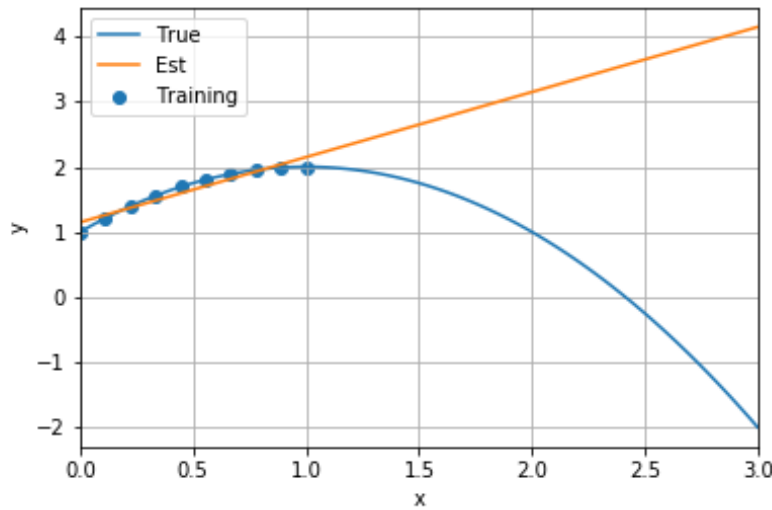


Figure 1: True function $f_0(x)$, estimate $f(x, \hat{\beta})$ and training points (x_i, y_i)

```
plt.ylabel('y')
plt.savefig('bias.png')
```

The resulting figure is shown in Fig. 1.

- (d) We see that the linear fit attempts to fit the quadratic in the region $x \in [0, 1]$ where the training data was. But, the fit is poor outside this region. In particular, in the interval $[0, 3]$, the bias error (difference between the true and estimated function) is largest at $x = 3$.
5. A medical researcher wishes to evaluate a new diagnostic test for cancer. A clinical trial is conducted where the diagnostic measurement y of each patient is recorded along with attributes of a sample of cancerous tissue from the patient. Three possible models are considered for the diagnostic measurement:
- Model 1: The diagnostic measurement y depends linearly only on the cancer volume.
 - Model 2: The diagnostic measurement y depends linearly on the cancer volume and the patient's age.
 - Model 3: The diagnostic measurement y depends linearly on the cancer volume and the patient's age, but the dependence (slope) on the cancer volume is different for two types of cancer – Type I and II.
- (a) Define variables for the cancer volume, age and cancer type and write a linear model for the predicted value \hat{y} in terms of these variables for each of the three models above. For Model 3, you will want to use one-hot coding.
- (b) What are the numbers of parameters in each model? Which model is the most complex?

- (c) Since the models in part (a) are linear, given training data, we should have $\hat{\mathbf{y}} = \mathbf{A}\boldsymbol{\beta}$ where $\hat{\mathbf{y}}$ is the vector of predicted values on the training data, \mathbf{A} is a feature matrix and $\boldsymbol{\beta}$ is the vector of parameters. To test the different models, data is collected from 100 patients. The records of the first three patients are shown below:

Patient ID	Measurement y	Cancer type	Cancer volume	Patient age
12	5	I	0.7	55
34	10	II	1.3	65
23	15	II	1.6	70
\vdots	\vdots	\vdots	\vdots	\vdots

Based on this data, what would be the values of first three rows of the three \mathbf{A} matrices be for the three models in part (a)?

- (d) To evaluate the models, 10-fold cross validation is used with the following results.

Model	Mean training RSS	Mean test RSS	Test RSS std deviation
1	2.0	2.01	0.03
2	0.7	0.72	0.04
3	0.65	0.70	0.05

All RSS values are per sample, and the last column is the (biased) standard deviation – not the standard error. Which model should be selected based on the “one standard error rule”?

Solution

- (a) Let x_1 be the cancer volume, x_2 be the patient’s age and x_3 be the cancer type:

$$x_3 = \begin{cases} 0 & \text{cancer is Type I} \\ 1 & \text{cancer is Type II,} \end{cases}$$

Then, the models can be written as:

$$\text{Model 1: } \hat{y} = \beta_0 + \beta_1 x_1,$$

$$\text{Model 2: } \hat{y} = \beta_0 + \beta_1 x_1 + \beta_2 x_2,$$

$$\text{Model 3: } \hat{y} = \beta_0 + \beta_1 x_1(1 - x_3) + \beta_2 x_1 x_3 + \beta_3 x_2.$$

In Model 3, we have used one-hot coding on the slope of x_1 . Specifically, when $x_3 = 0$ (Type I cancer), the slope for x_1 is β_1 ; when $x_3 = 1$ (Type II cancer), the slope for x_1 is β_2 .

- (b) Models 1, 2 and 3 have 2, 3 and 4 parameters respectively. Model 3 is most complex.
(c) For Model 1, the first three rows of the feature matrix are:

$$\mathbf{A} = \begin{bmatrix} 1 & x_{11} \\ 1 & x_{21} \\ 1 & x_{31} \\ \vdots & \vdots \end{bmatrix} = \begin{bmatrix} 1 & 0.7 \\ 1 & 1.3 \\ 1 & 1.6 \\ \vdots & \vdots \end{bmatrix}.$$

For Model 2, the first three rows of the feature matrix are:

$$\mathbf{A} = \begin{bmatrix} 1 & x_{11} & x_{12} \\ 1 & x_{21} & x_{22} \\ 1 & x_{31} & x_{32} \\ \vdots & & \vdots \end{bmatrix} = \begin{bmatrix} 1 & 0.7 & 55 \\ 1 & 1.3 & 65 \\ 1 & 1.6 & 70 \\ \vdots & & \vdots \end{bmatrix}.$$

For Model 3, the first three rows of the feature matrix are:

$$\mathbf{A} = \begin{bmatrix} 1 & x_{11}(1 - x_{13}) & x_{11}x_{13} & x_{12} \\ 1 & x_{21}(1 - x_{23}) & x_{21}x_{23} & x_{22} \\ 1 & x_{31}(1 - x_{33}) & x_{31}x_{33} & x_{32} \\ \vdots & & \vdots & \end{bmatrix} = \begin{bmatrix} 1 & 0.7 & 0 & 55 \\ 1 & 0 & 1.3 & 65 \\ 1 & 0 & 1.6 & 70 \\ \vdots & & \vdots & \end{bmatrix}.$$

- (d) The lowest test error is for Model 3 with a mean RSS = 0.70. The standard deviation is 0.05, so the standard error is

$$SE = 0.05/\sqrt{K-1} = 0.05/\sqrt{9} = 0.0167.$$

Note the use of $\sqrt{K-1}$ since the standard deviation was biased. Hence, the RSS target is $0.70 + 0.0167 = 0.7167$. The least complex model below this target is Model 3.